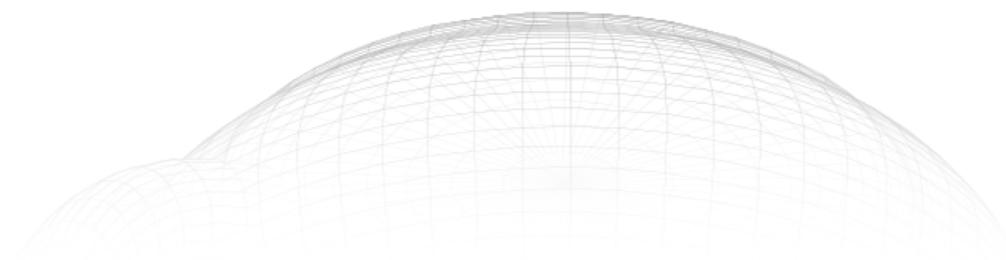
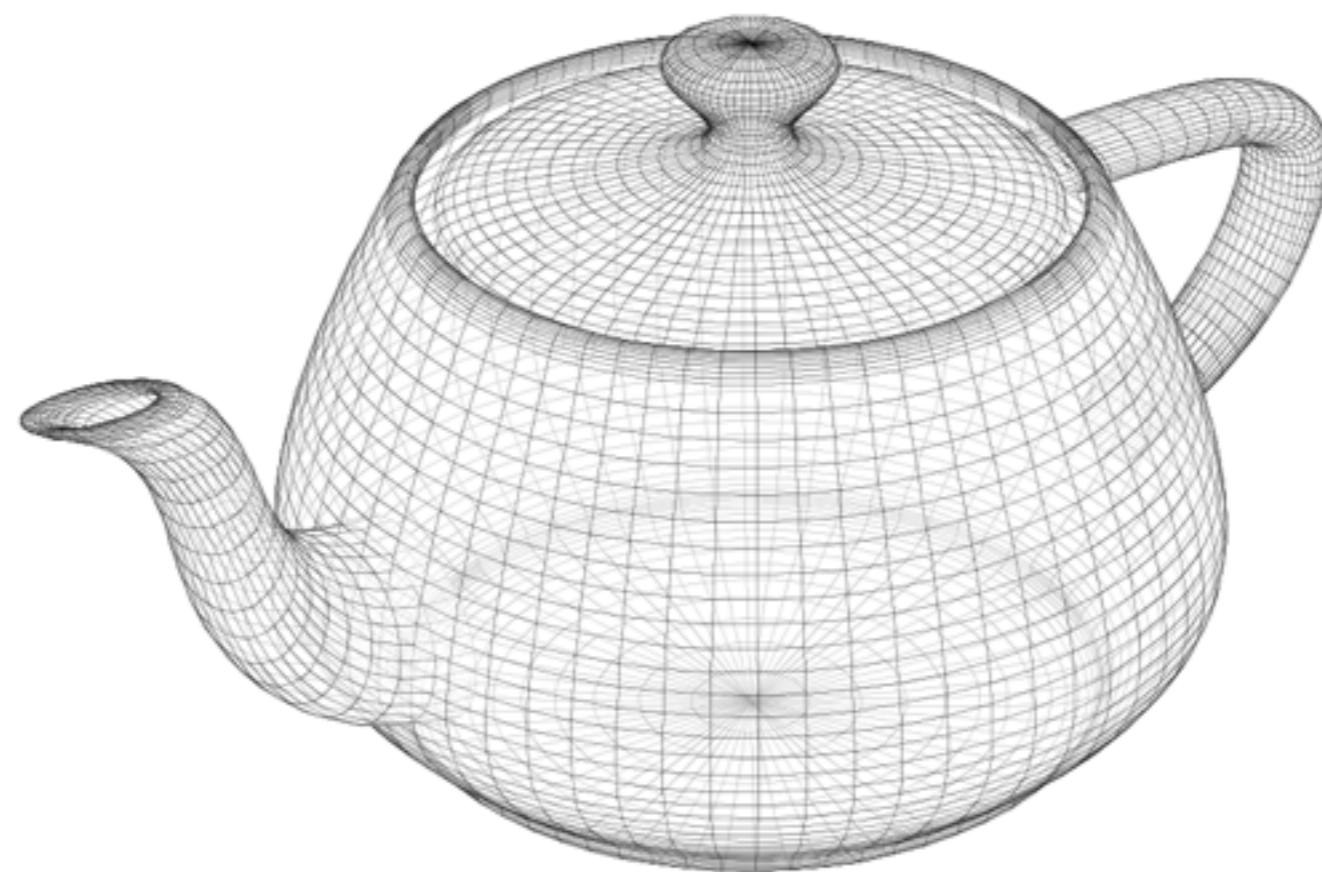


INTRODUCTION TO GPU COMPUTING

Ilya Kuzovkin

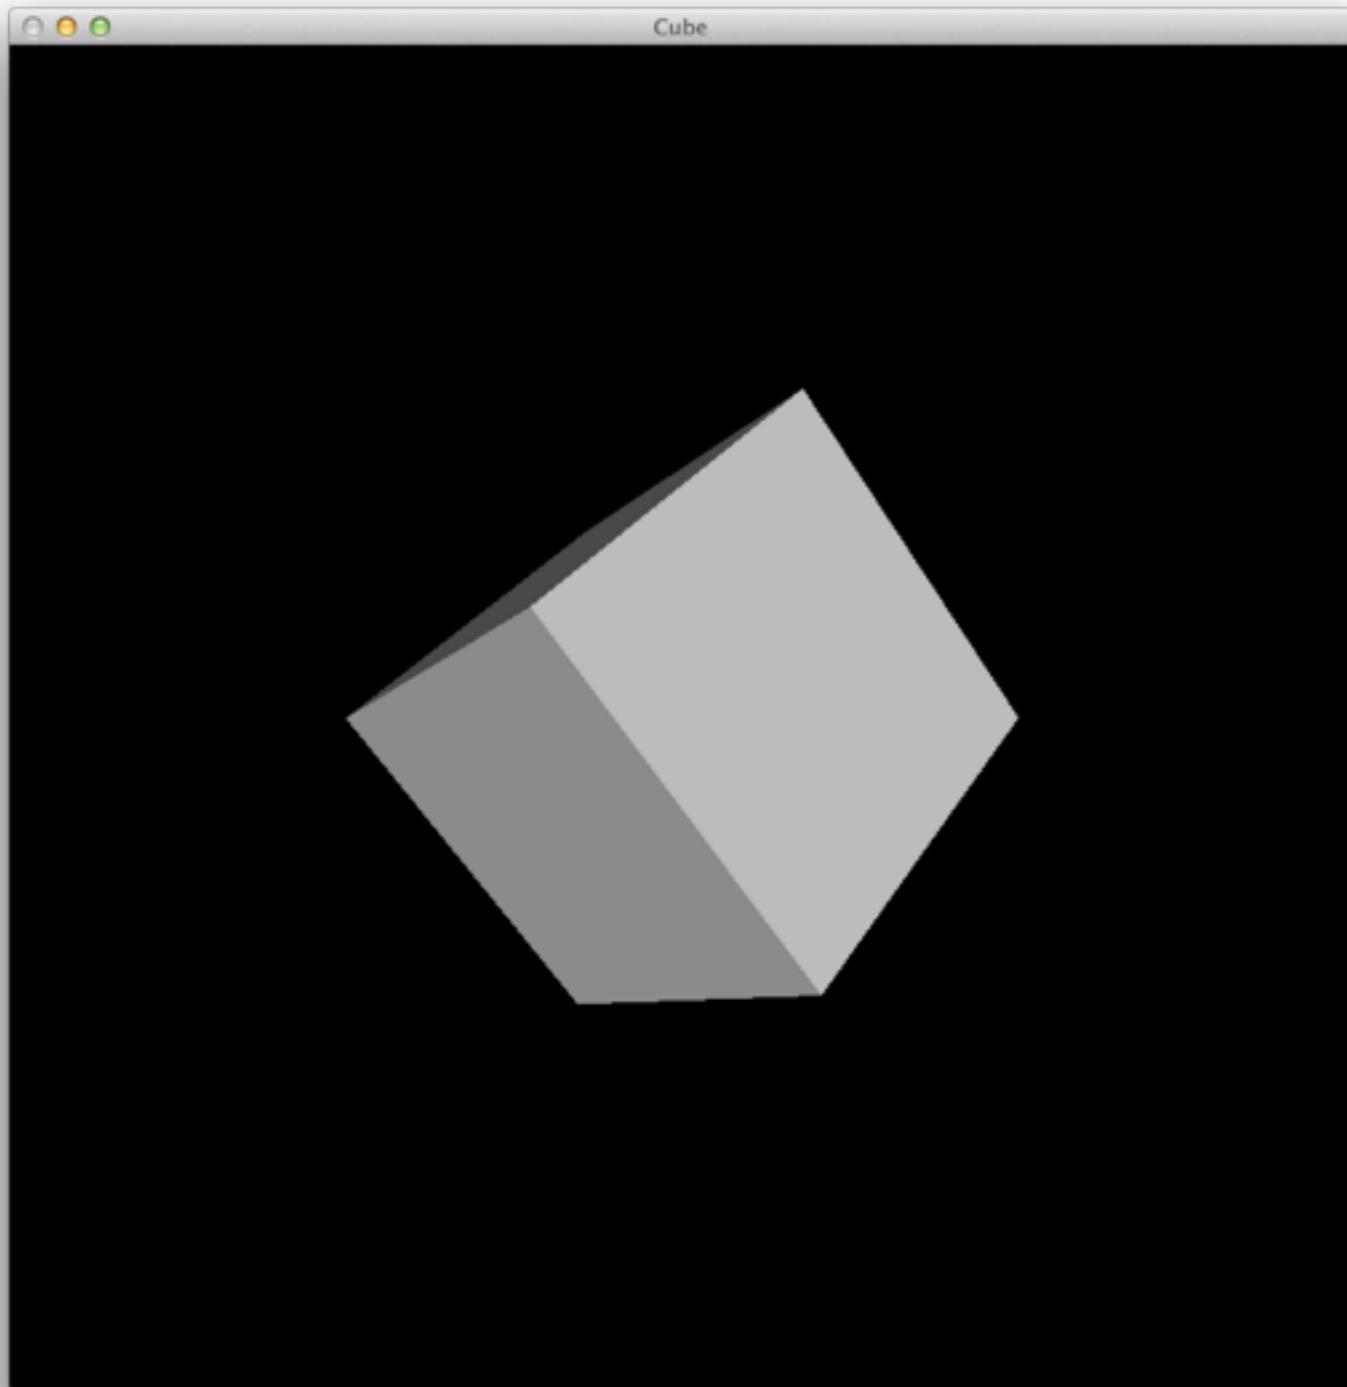


13 May 2014, Tartu



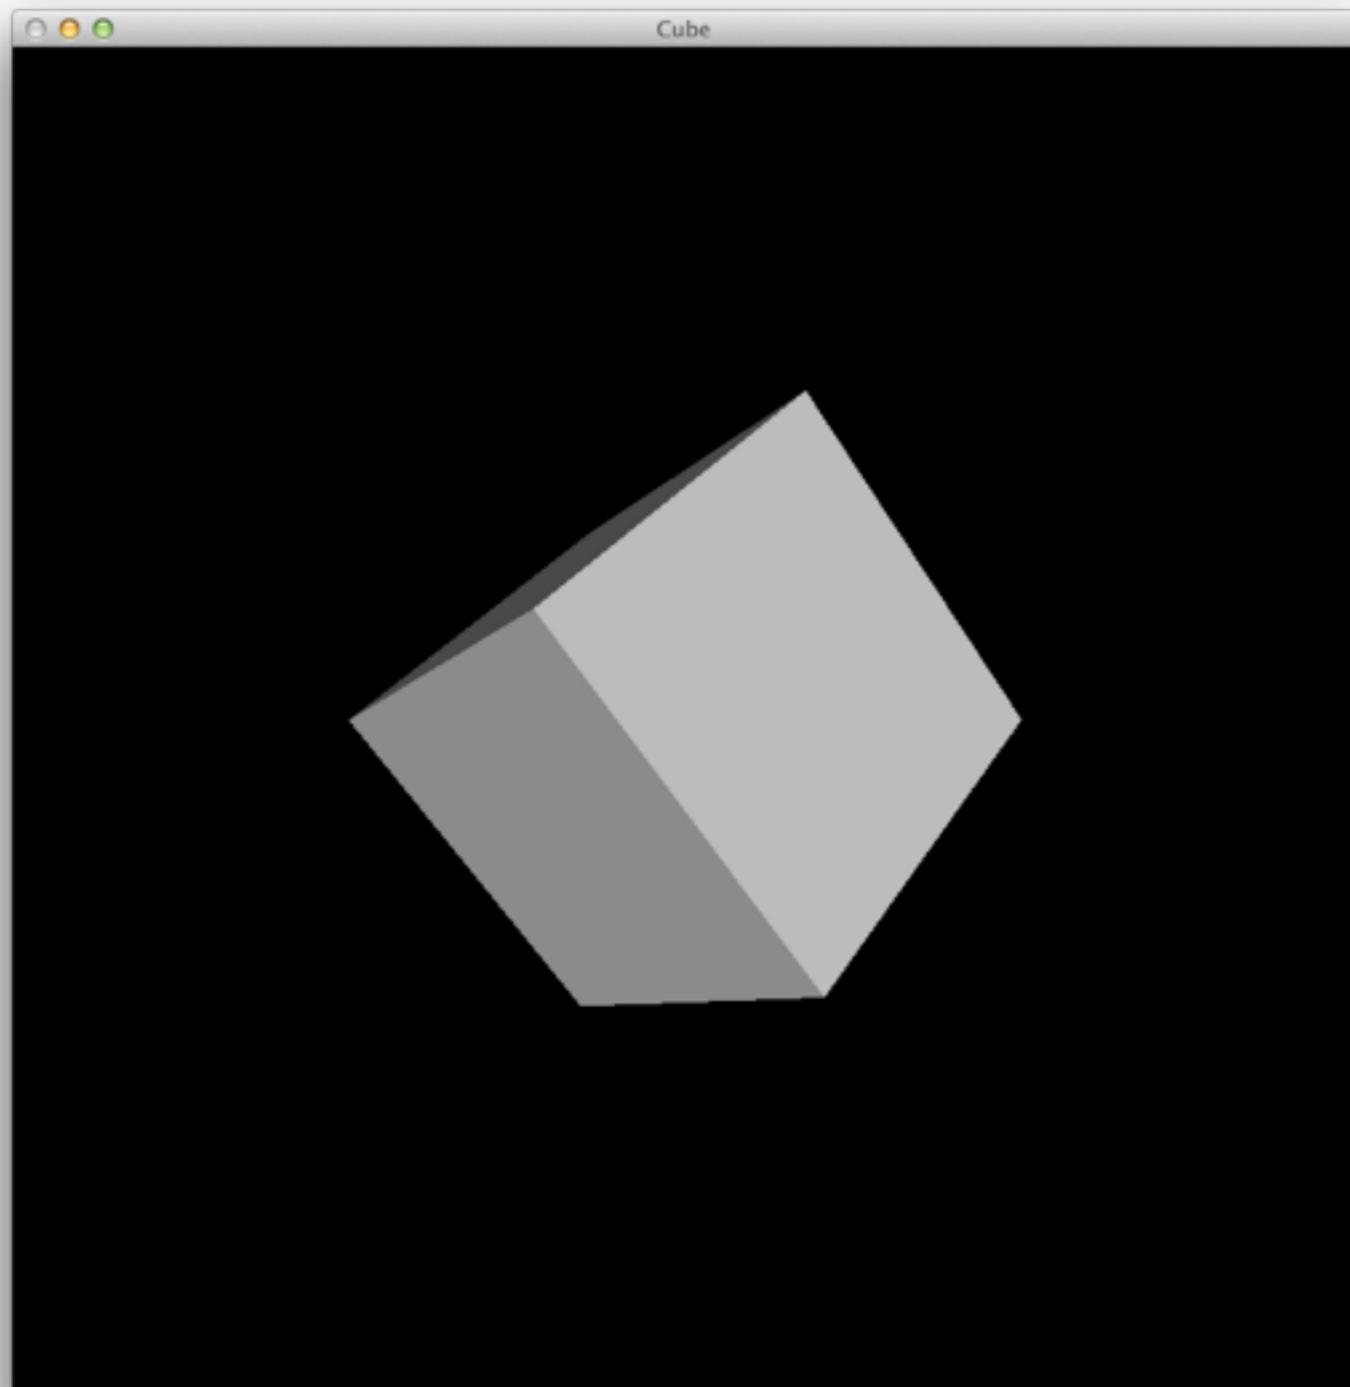
PART I
“TEAPOT”

SIMPLE OPENGL PROGRAM



Idea of computing on GPU emerged because GPUs became very good at parallel computations.

SIMPLE OPENGL PROGRAM



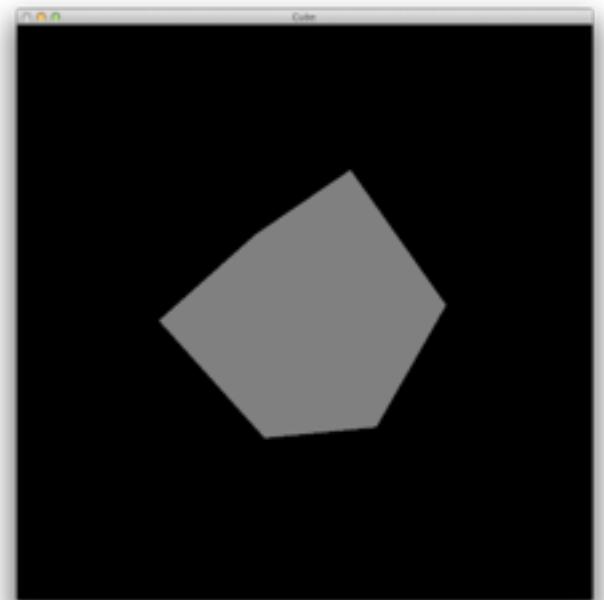
Idea of computing on GPU emerged because GPUs became very good at parallel computations.

Let us start from observing an example of parallelism in a simple OpenGL application.

SIMPLE OPENGL PROGRAM

You will need CodeBlocks^{Windows, Linux} or XCode^{Mac} to run this example.

- Install CodeBlocks bundled with MinGW compiler from <http://www.codeblocks.org/downloads/26>
- Download codebase from <https://github.com/kuz/Introduction-to-GPU-Computing>
- Open the project from the **code/Cube**
- Compile & run it



SHADER PROGRAM

Program which is **executed** on GPU.

Has to be written using **shading language**.

In OpenGL this language is **GLSL**, which is based on C.

SHADER PROGRAM

Program which is **executed** on GPU.

Has to be written using **shading language**.

In OpenGL this language is **GLSL**, which is based on C.

OpenGL has 5 main shader stages:

- Vertex Shader
- Tessellation Control
- Geometry Shader
- Fragment Shader
- Compute Shader (since 4.3)

SHADER PROGRAM

Program which is **executed** on GPU.

Has to be written using **shading language**.

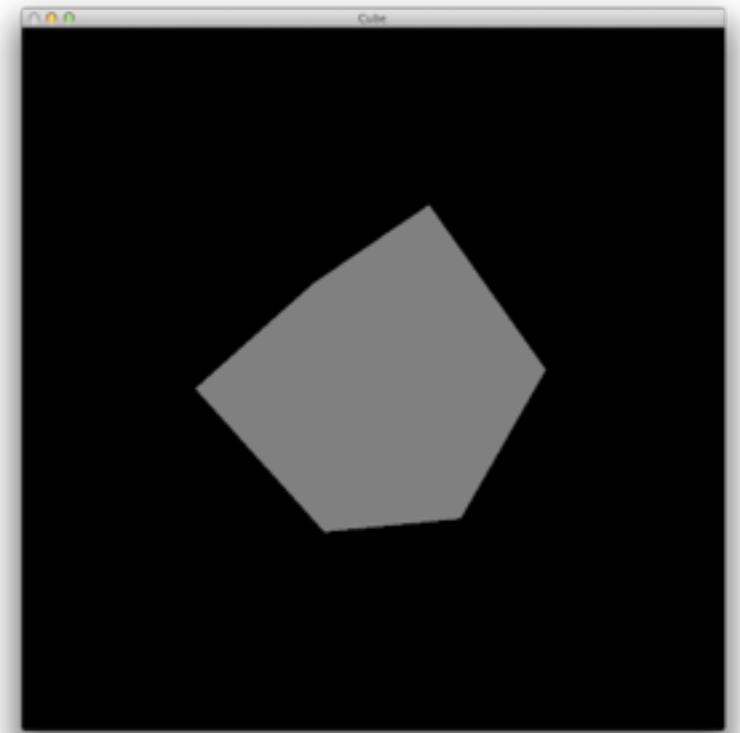
In OpenGL this language is **GLSL**, which is based on C.

OpenGL has 5 main shader stages:

- Vertex Shader
- Tessellation Control
- Geometry Shader
- Fragment Shader
- Compute Shader (since 4.3)

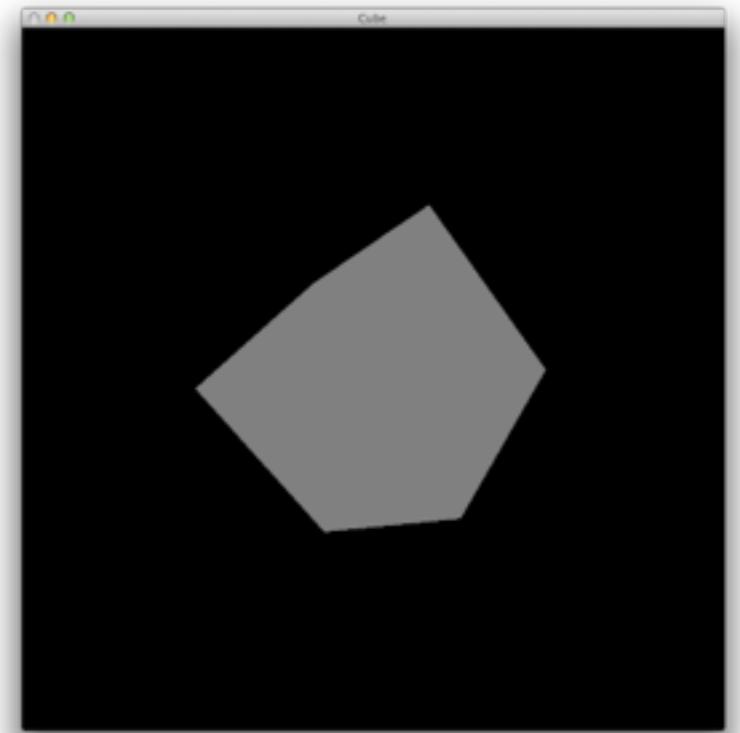
LIGHTING

Is it a cube or not?
We will find out as soon as
we add **lighting** to the scene.



LIGHTING

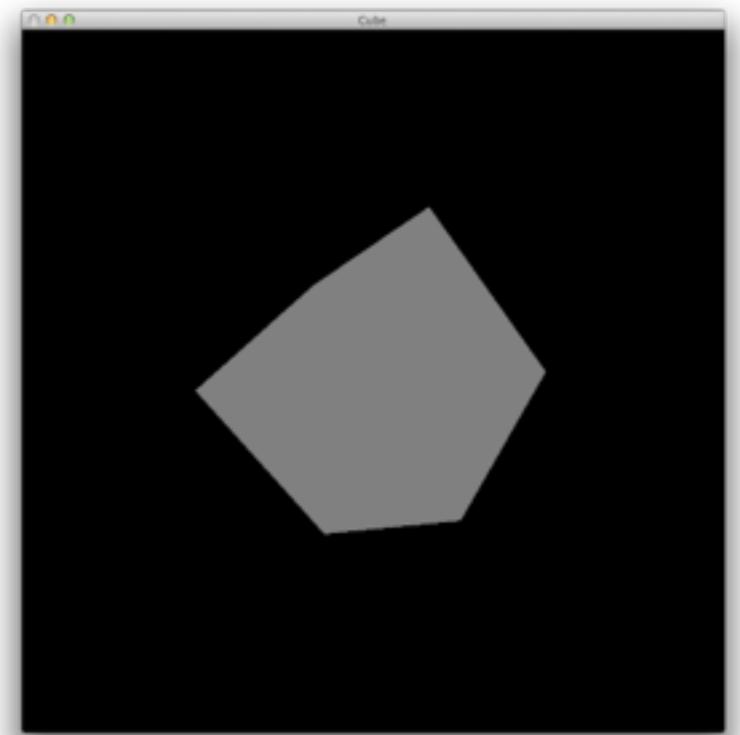
Is it a cube or not?
We will find out as soon as
we add **lighting** to the scene.



$$I_A \cdot A + I_D \cdot D \cdot \mathbf{l}^T \mathbf{n} + I_S \cdot S \cdot (\mathbf{r}^T \mathbf{v})^s$$

LIGHTING

Is it a cube or not?
We will find out as soon as
we add **lighting** to the scene.



$$I_A \cdot A + I_D \cdot D \cdot \mathbf{l}^T \mathbf{n} + I_S \cdot S \cdot (\mathbf{r}^T \mathbf{v})^s$$

Exercise: code that equation into
fragment shader of the Cube program

LIGHTING

```
// Configuration
vec3 color = vec3(0, 0, 0);
vec3 light = vec3(0, 0, 1);
vec3 viewer = vec3(0, 0, 1);
vec3 normal = vertexNormal;

// Ambient
vec3 l_ambient = vec3(0.2, 0.2, 0.2);
vec3 m_ambient = vec3(0.5, 0.5, 0.5);
color += l_ambient * m_ambient;

// Diffuse
vec3 l_diffuse = vec3(0.7, 0.7, 0.7);
vec3 m_diffuse = vec3(0.5, 0.5, 0.5);
color += l_diffuse * m_diffuse * dot(light, normal);

// Specular
vec3 l_specular = vec3(1.0, 1.0, 1.0);
vec3 m_specular = vec3(0.5, 0.5, 0.5);
float sh = 20.0;
vec3 reflection = reflect(-light, normal);
color += l_specular * m_specular * pow(clamp(dot(reflection, viewer), 0, 1), sh);

// Return final color
fragColor = vec4(color, 0);
```

COMPARE FPS

- Run the program with **lighting enabled** and look at **FPS** values

COMPARE FPS

- Run the program with **lighting enabled** and look at **FPS** values
- In **cube.cpp** **idle()** function uncomment **dummy** code which simulates approximately **same** amount of computations as Phong lighting model requires.

COMPARE FPS

- Run the program with **lighting enabled** and look at **FPS** values
- In `cube.cpp` `idle()` function uncomment **dummy** code which simulates approximately **same** amount of computations as Phong lighting model requires.
- Note that these computations are performed on **CPU**

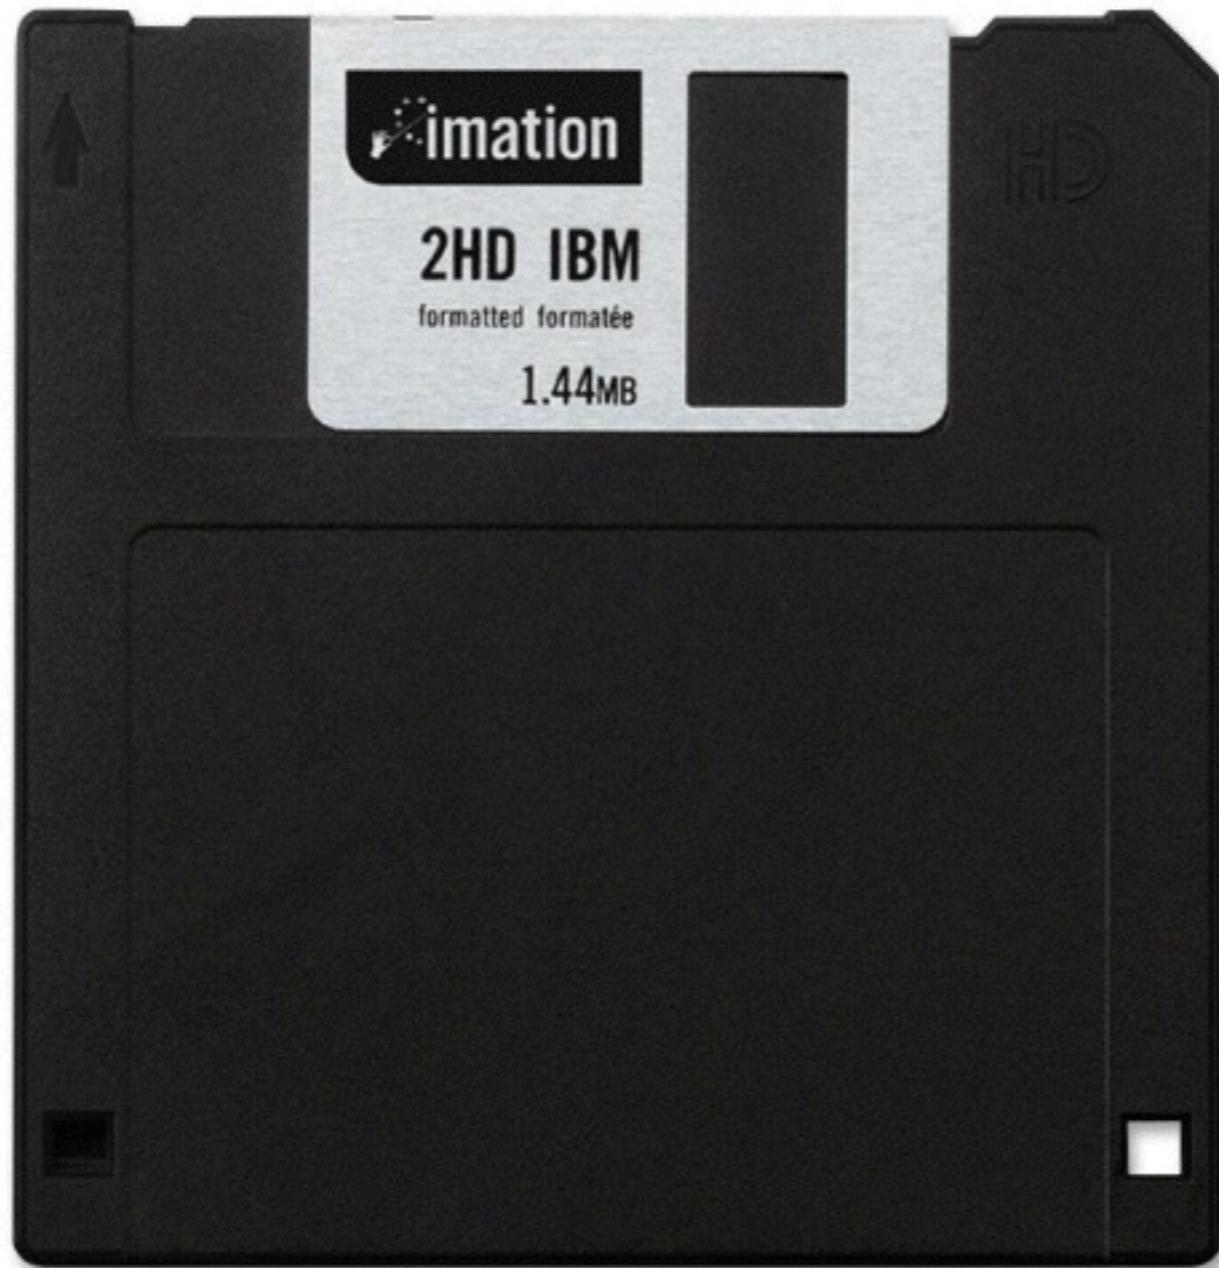
COMPARE FPS

- Run the program with **lighting enabled** and look at **FPS** values
- In `cube.cpp` `idle()` function uncomment **dummy** code which simulates approximately **same** amount of computations as Phong lighting model requires.
- Note that these computations are performed on **CPU**
- Observe how **FPS** has **changed**

COMPARE FPS

- Run the program with **lighting enabled** and look at **FPS** values

- Parallel computations are fast on GPU.
Let's use it to compute something useful.
- Note that those computations are performed on GPU
- Observe how FPS has changed



PART II “OLD SCHOOL”

OPENGL PIPELINE + GLSL

Take the **input data** from
the CPU memory and put
it as an **image** into the
GPU memory

OPENGL PIPELINE + GLSL

Take the **input data** from the CPU memory and put it as an **image** into the GPU memory



In the fragment shader perform a computation on each of the pixels of that **image**

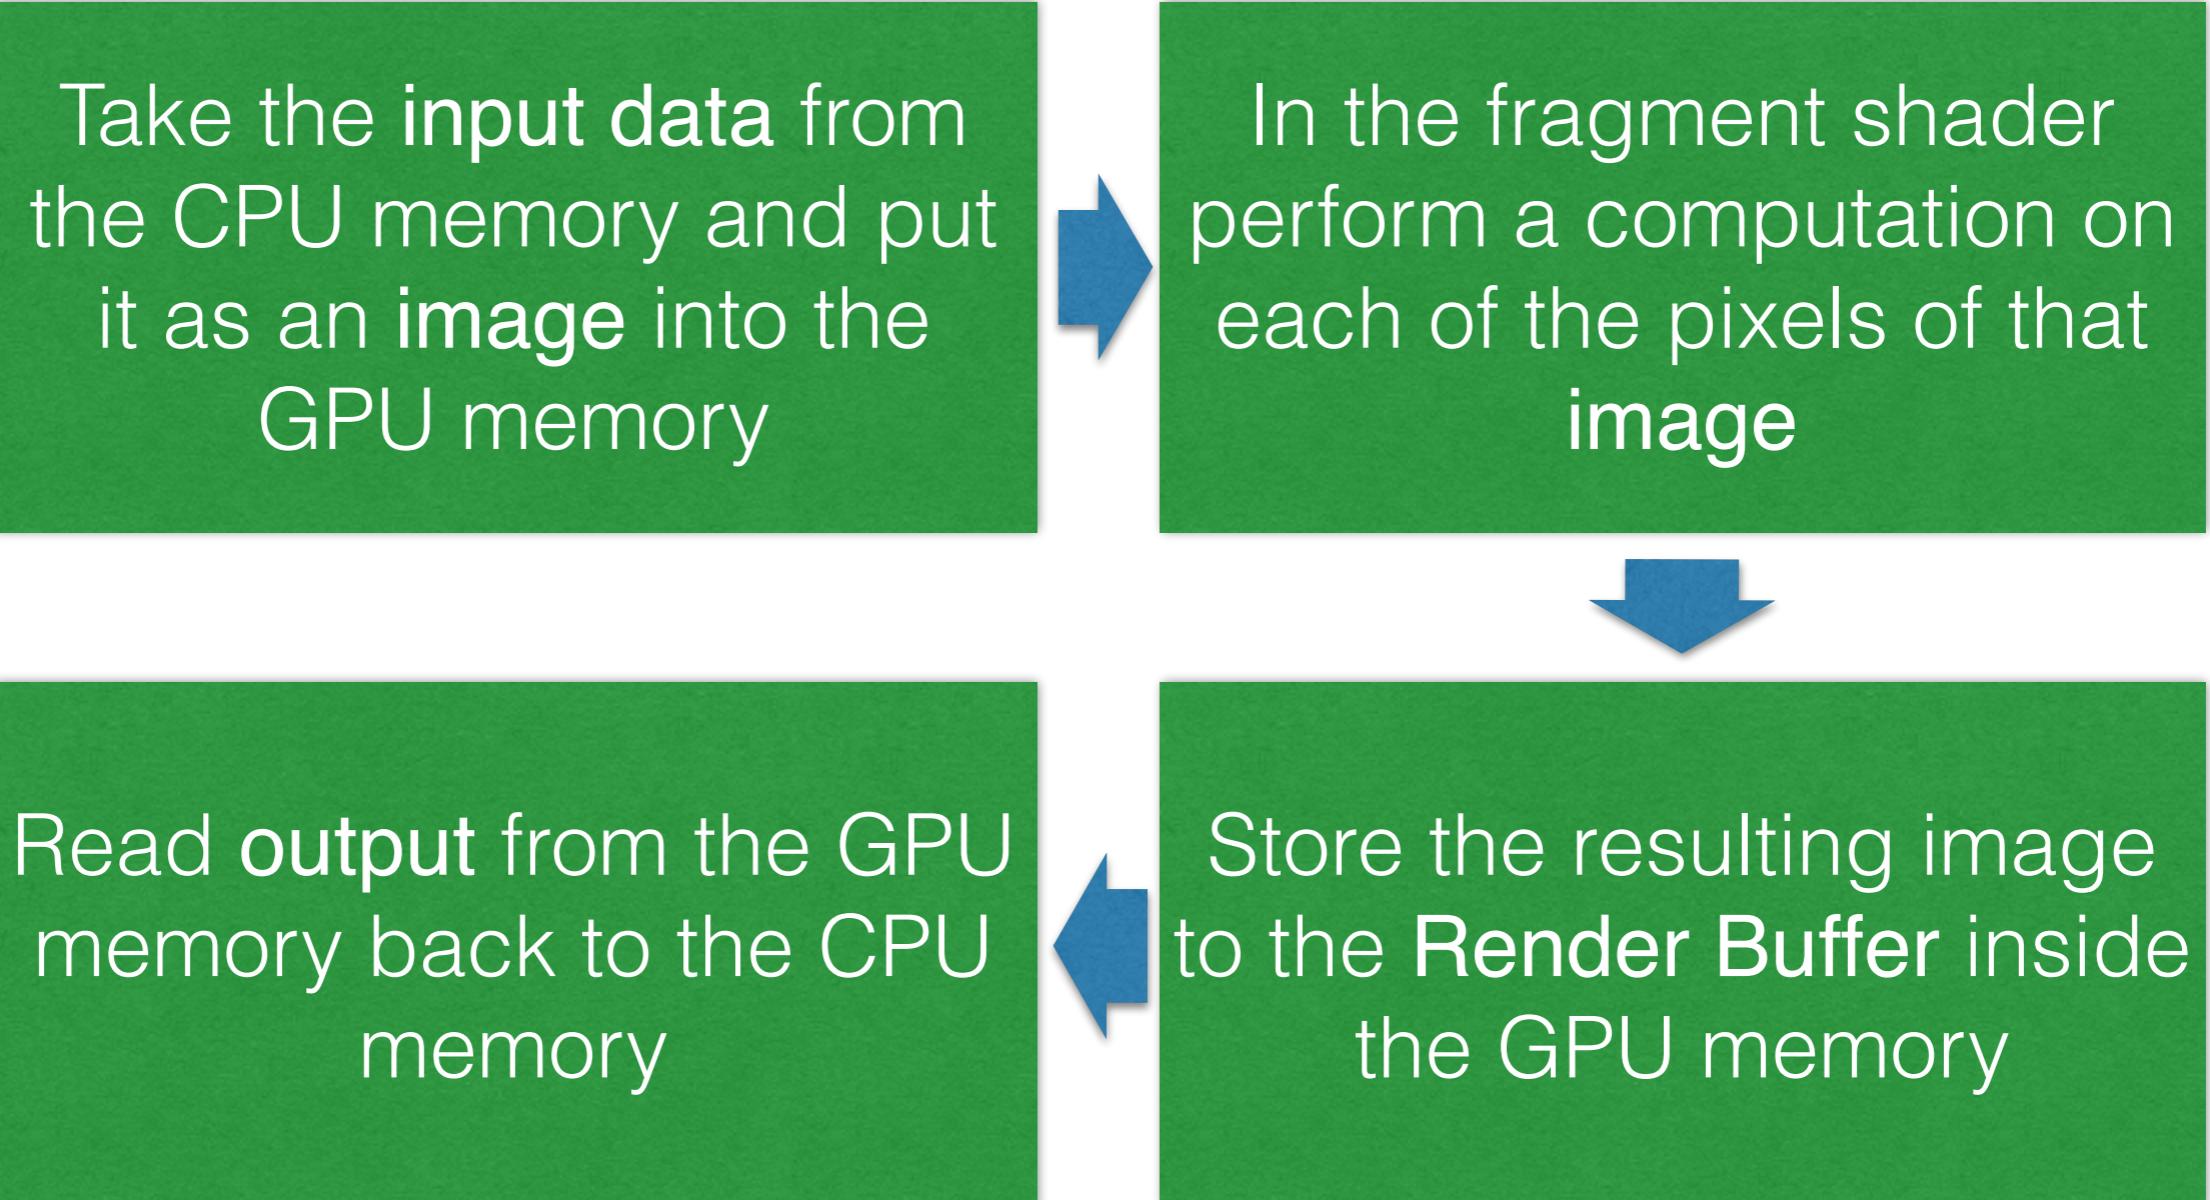
OPENGL PIPELINE + GLSL

Take the **input data** from the CPU memory and put it as an **image** into the GPU memory

In the fragment shader perform a computation on each of the pixels of that **image**

Store the resulting image to the **Render Buffer** inside the GPU memory

OPENGL PIPELINE + GLSL



OPENGL PIPELINE + GLSL

- Create **texture** where will store the **input data**

```
glGenTextures(1, &fbo->texid);
glBindTexture(GL_TEXTURE_2D, fbo->texid);
glTexImage2D(GL_TEXTURE_2D, 0, GL_RGBA32F_ARB,
             width, height, 0, GL_RGBA, GL_FLOAT, data);
```

OPENGL PIPELINE + GLSL

- Create **texture** where will store the **input data**

```
glGenTextures(1, &fbo->texid);
glBindTexture(GL_TEXTURE_2D, fbo->texid);
glTexImage2D(GL_TEXTURE_2D, 0, GL_RGBA32F_ARB,
             width, height, 0, GL_RGBA, GL_FLOAT, data);
```

- Create **FrameBuffer Object (FBO)** to “render” to

```
glGenFramebuffers(1, &fbo->fb);           // frame buffer id
glGenRenderbuffers(1, &fbo->rb);           // render buffer id
glBindFramebuffer(GL_FRAMEBUFFER, fbo->fb);
glFramebufferTexture2D(GL_FRAMEBUFFER, GL_COLOR_ATTACHMENT0,
                      GL_TEXTURE_2D, fbo->texid, 0);
```

OPENGL PIPELINE + GLSL

- Run OpenGL pipeline

OPENGL PIPELINE + GLSL

- Run OpenGL pipeline
 - Render GL_QUADS of same size as the **texture** matrix

OPENGL PIPELINE + GLSL

- Run OpenGL pipeline
 - Render GL_QUADS of same size as the **texture** matrix
 - Use **fragment shader** to perform per-fragment computations using data from the **texture**

OPENGL PIPELINE + GLSL

- Run OpenGL pipeline
 - Render GL_QUADS of same size as the **texture** matrix
 - Use **fragment shader** to perform per-fragment computations using data from the **texture**
 - OpenGL will **store result** in the texture given to the **Render Buffer** (within Framebuffer Object)

OPENGL PIPELINE + GLSL

- Run OpenGL pipeline
 - Render GL_QUADS of same size as the **texture** matrix
 - Use **fragment shader** to perform per-fragment computations using data from the **texture**
 - OpenGL will **store result** in the texture given to the **Render Buffer** (within Framebuffer Object)
- Read the data from the Render Buffer

```
glReadBuffer(GL_COLOR_ATTACHMENT0);
glReadPixels(0, 0, texSize, texSize,
             GL_RGBA, GL_FLOAT, result);
```

OPENGL PIPELINE + GLSL

- Run OpenGL pipeline
 - Render GL_QUADS of same size as the **texture** matrix
 - Use **fragment shader** to perform per-fragment computations using data from the **texture**
 - OpenGL will **store result** in the texture given to the **Render Buffer** (within Framebuffer Object)

- Read the data from the Render Buffer

```
glReadBuffer(GL_COLOR_ATTACHMENT0);
glReadPixels(0, 0, texSize, texSize,
             GL_RGBA, GL_FLOAT, result);
```

- Can we use that to properly **debug GLSL**?

DEMO

The screenshot shows the Xcode IDE interface with the following details:

- Project Navigator:** Shows the project structure with targets: "NewTriangle" and "My Mac 64-bit". Frameworks included are GLUT.framework, libglfw.3.0.dylib, libGLEW.1.10.0.dylib, and OpenGL.framework.
- File Navigator:** Shows files: fbo.cpp, shader_util.cpp, fbo.vert.gls, and fbo.frag.gls.
- Code Editor:** The active file is fbo.cpp, displaying C++ code for initializing an FBO and rendering to it. The code uses OpenGL functions like glGenTextures, glBindTexture, glPixelStorei, glTexImage2D, glTexParameter, glGenFramebuffers, and glBindFramebuffer.
- Output Window:** Shows the terminal output of the application run. It includes:
 - Numerical values: 3.316625, 3.464102, 3.605551, 3.741657, 3.872983, 4.000000.
 - Timing information: Total ms (GPU): 225, Total ms (CPU): 2562.
 - Final message: Program ended with exit code: 0.

Run the project from the code/FBO



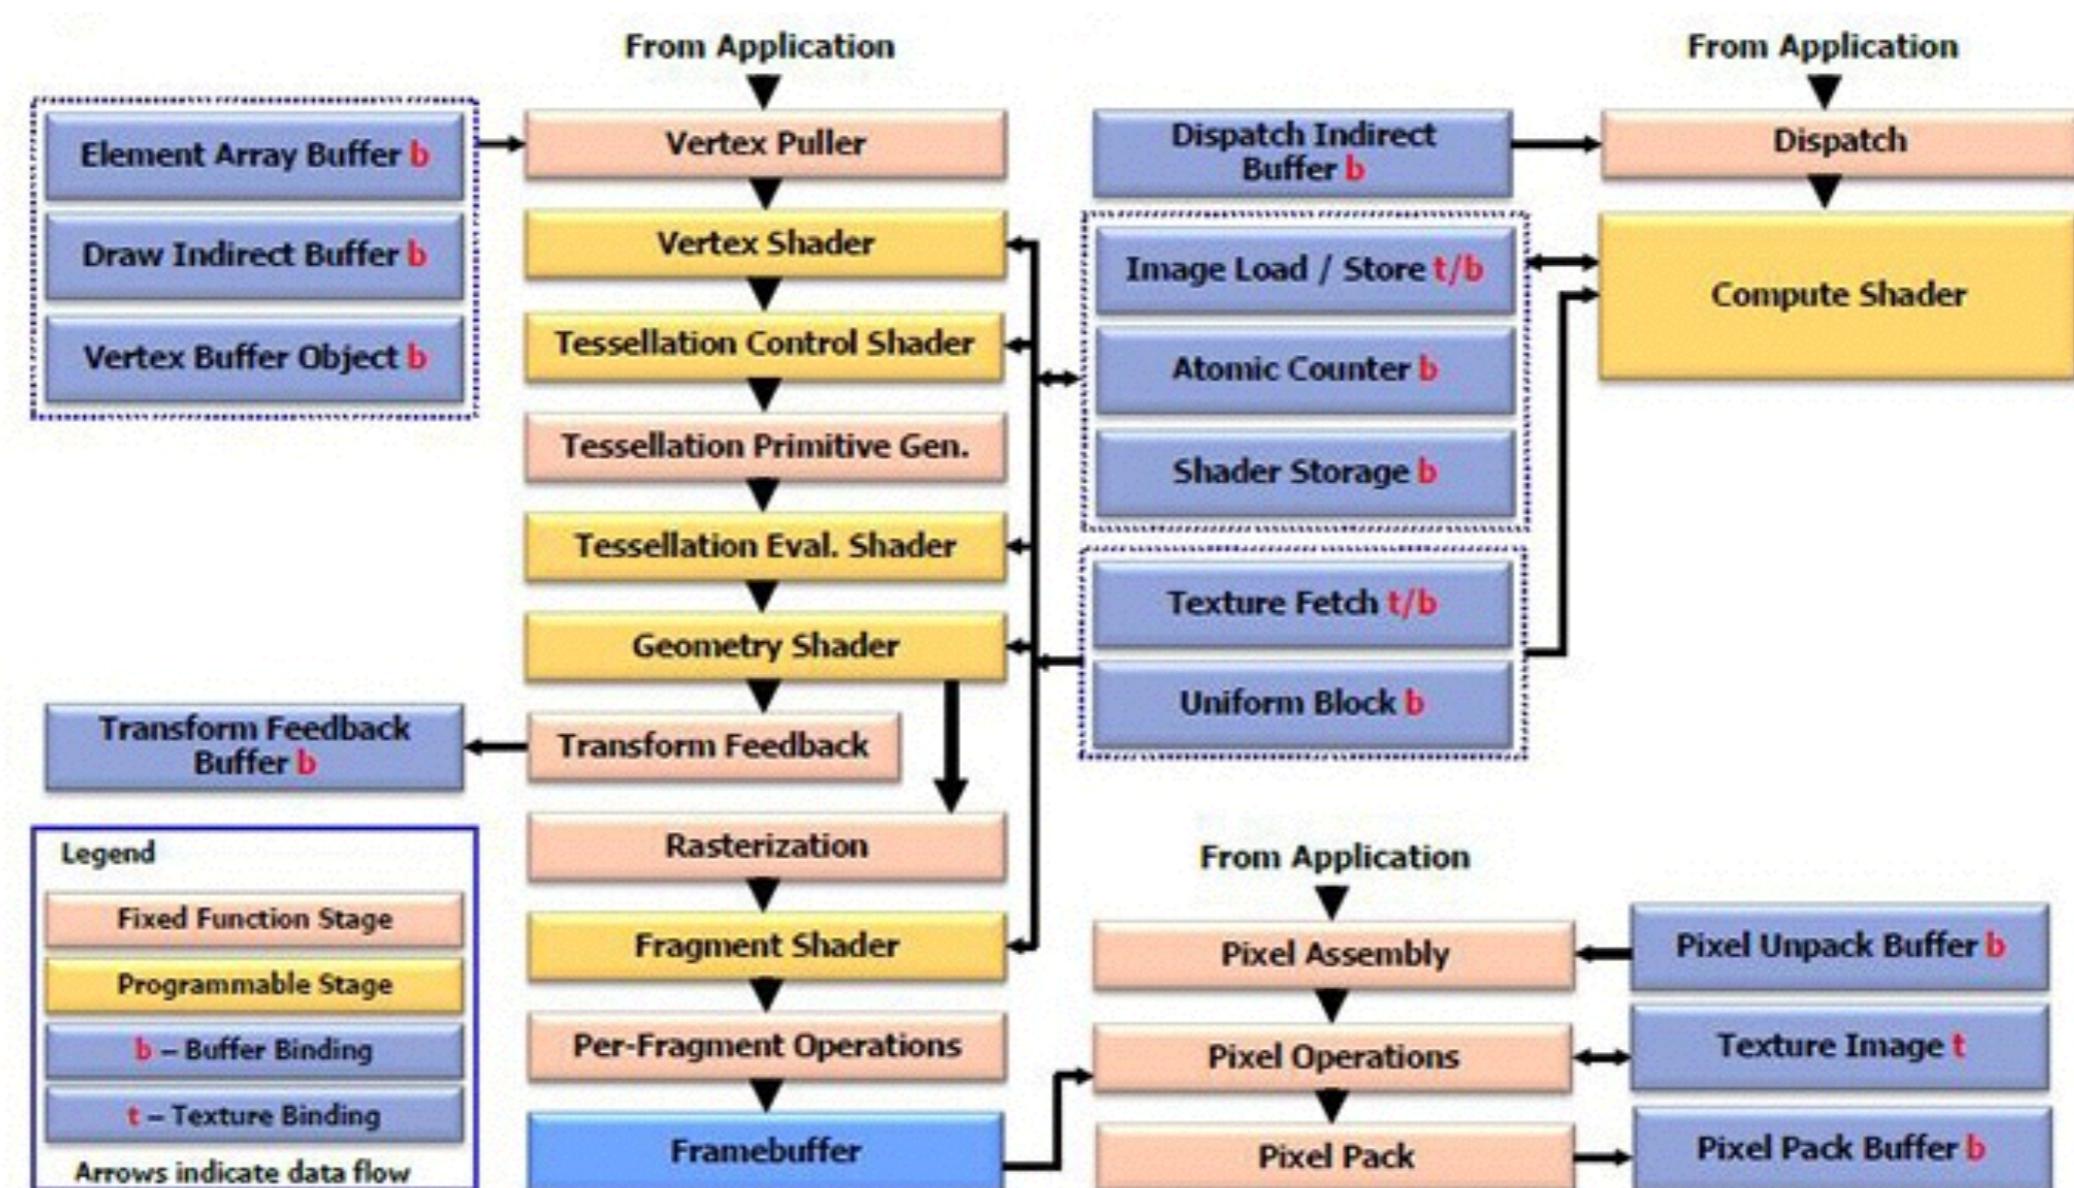
PART III
“MODERN TIMES”

COMPUTE SHADER

- Since OpenGL 4.3
- Used to compute things not related to rendering directly

COMPUTE SHADER

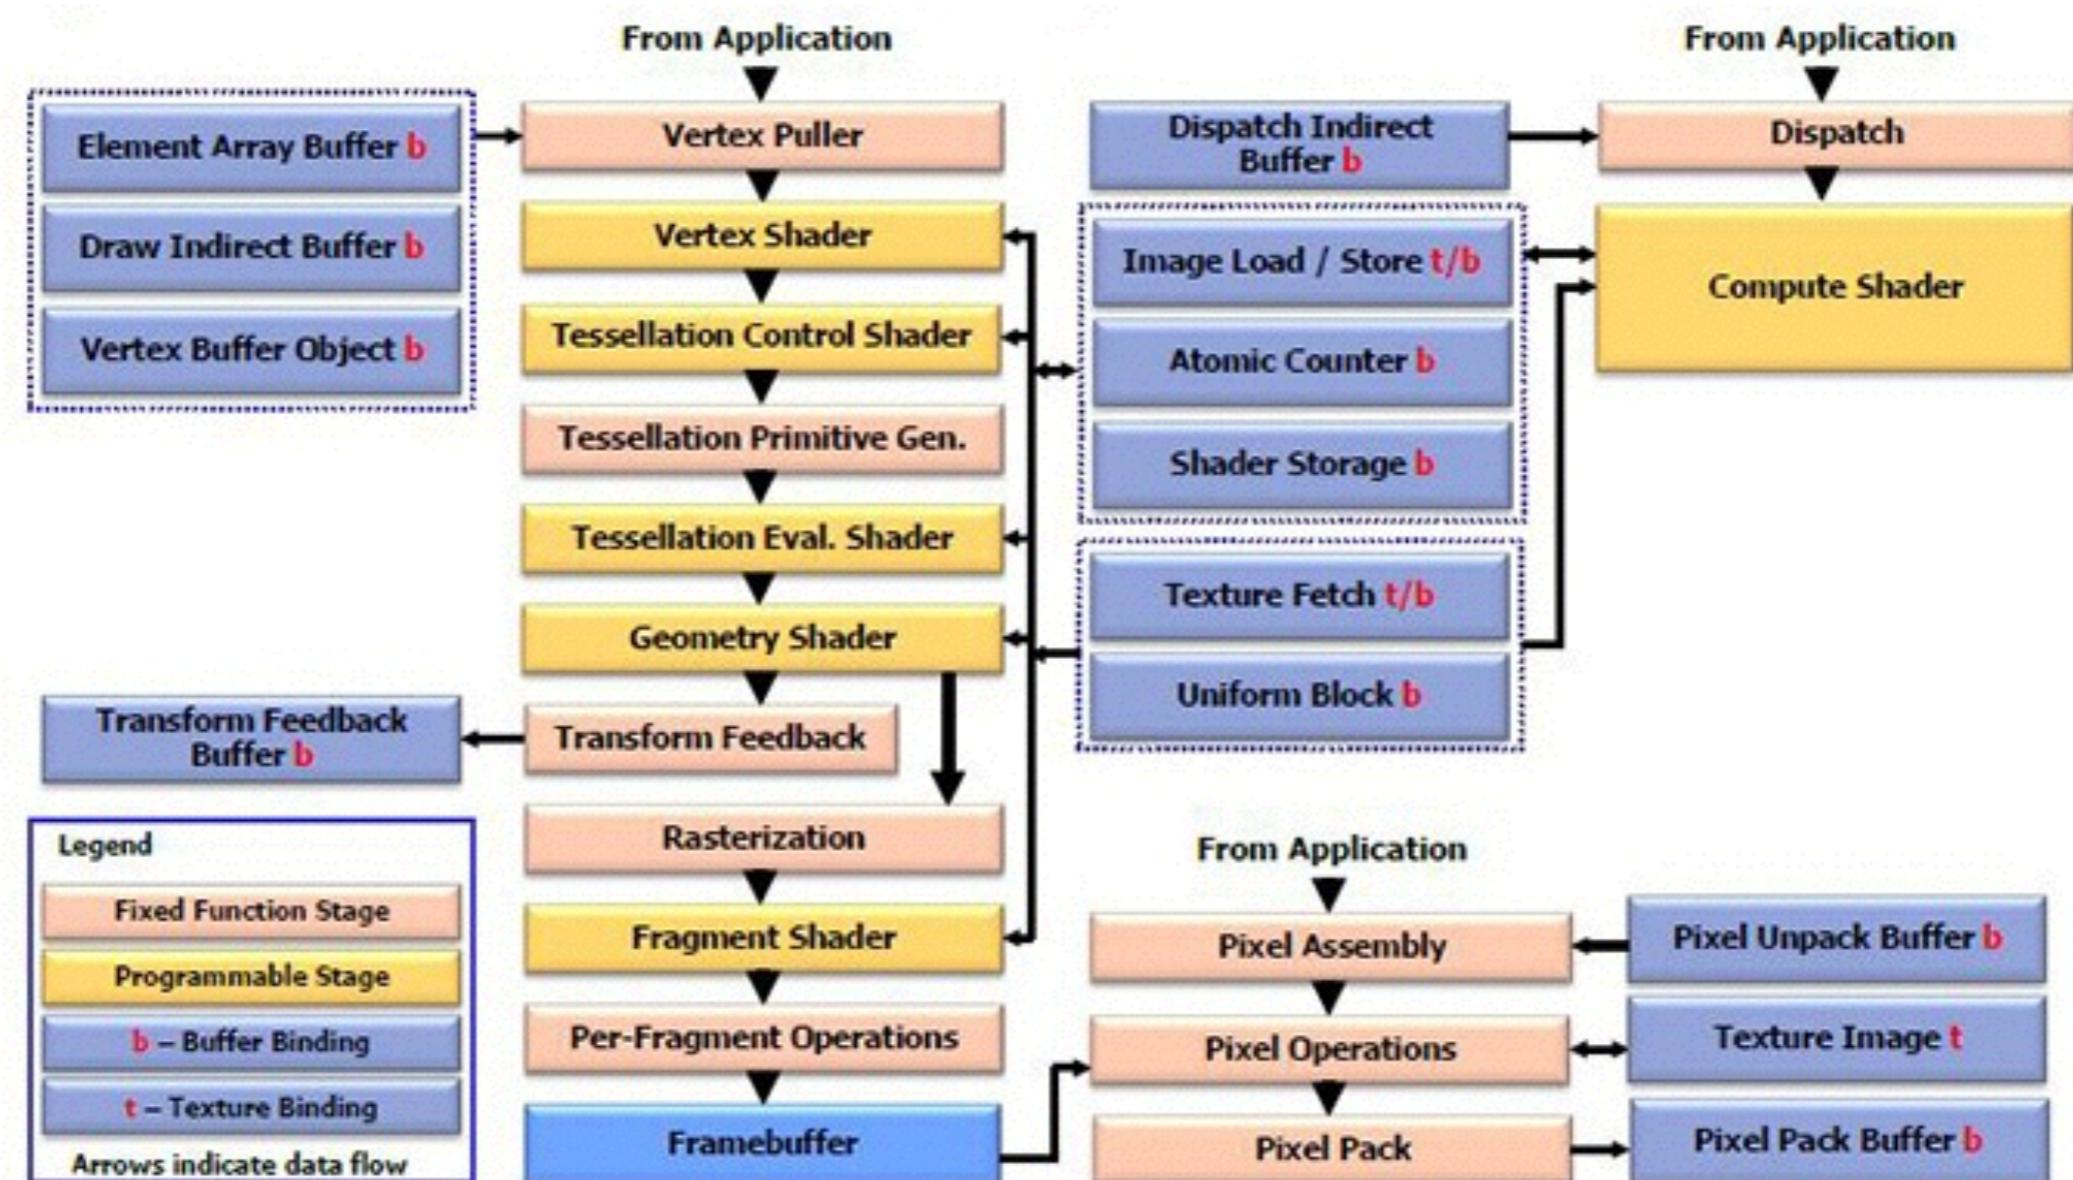
- Since OpenGL 4.3
- Used to compute things not related to rendering directly

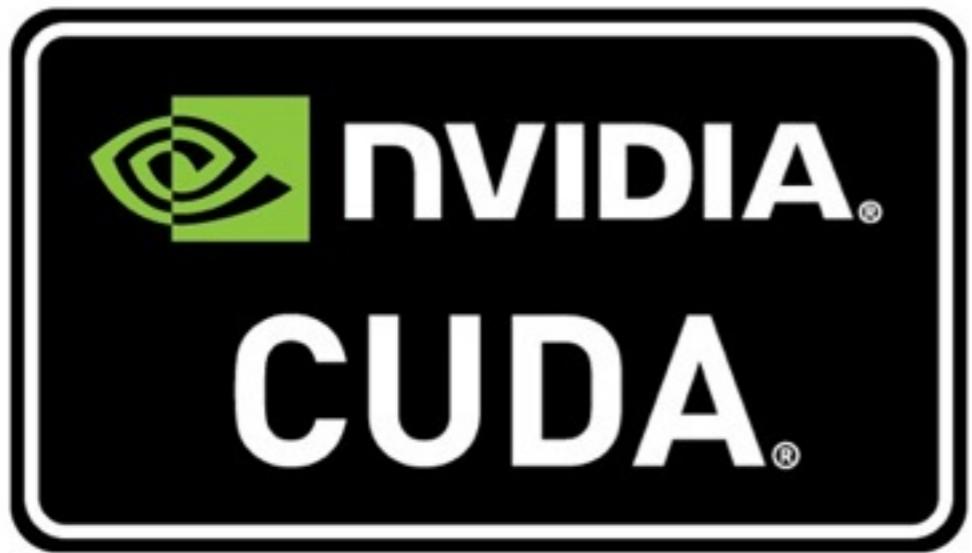


Will not talk
about it

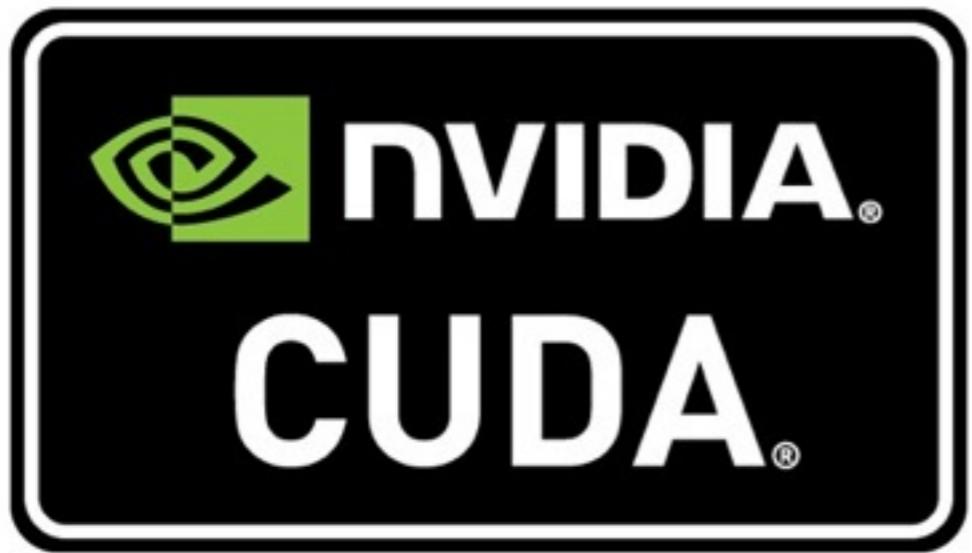
COMPUTE SHADER

- Since OpenGL 4.3
- Used to compute things not related to rendering directly





OpenCL



Supported only by
nVidia hardware

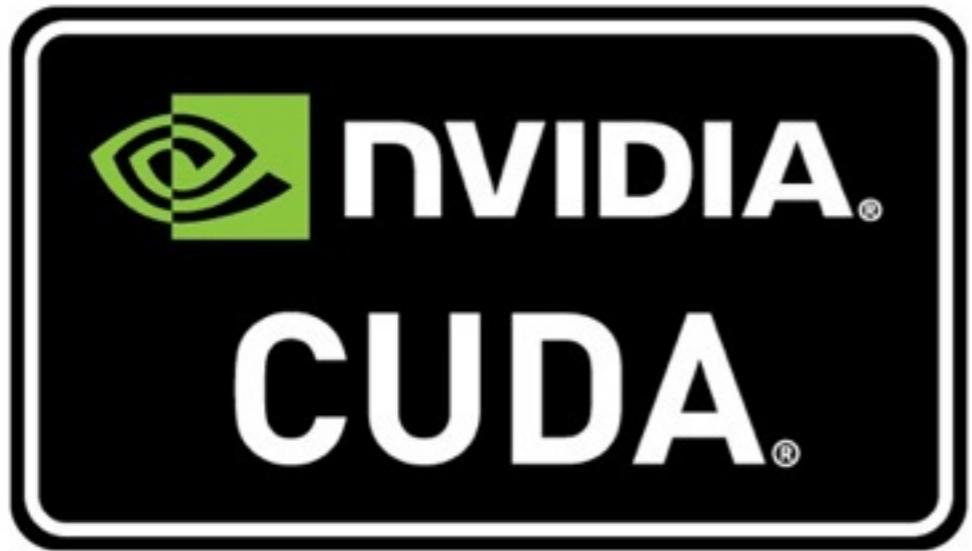
<https://developer.nvidia.com/cuda-gpus>



OpenCL

Supported by nVidia,
AMD, Intel, Qualcomm

<http://www.khronos.org/conformance/adopters/conformant-products#opencl>



Supported only by
nVidia hardware

<https://developer.nvidia.com/cuda-gpus>



OpenCL

Supported by nVidia,
AMD, Intel, Qualcomm

<http://www.khronos.org/conformance/adopters/conformant-products#opencl>

Implementations only
by nVidia

OpenCL



Supported only by
nVidia hardware

<https://developer.nvidia.com/cuda-gpus>



OpenCL

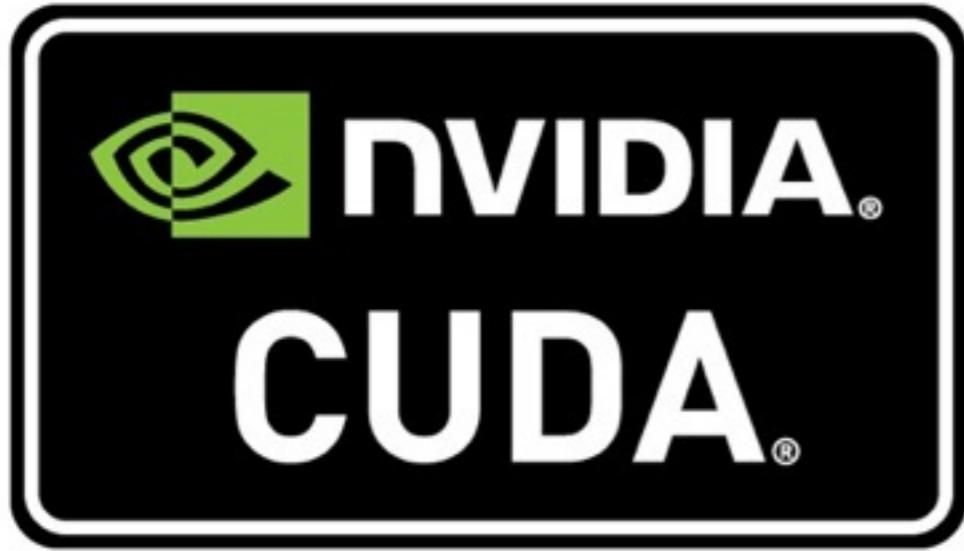
Supported by nVidia,
AMD, Intel, Qualcomm

<http://www.khronos.org/conformance/adopters/conformant-products#opencl>

Implementations only
by nVidia

OpenCL

~same performance levels



Supported only by
nVidia hardware

<https://developer.nvidia.com/cuda-gpus>



OpenCL

Supported by nVidia,
AMD, Intel, Qualcomm

<http://www.khronos.org/conformance/adopters/conformant-products#opencl>

Implementations only
by nVidia

OpenCL

~same performance levels

Developer-friendly

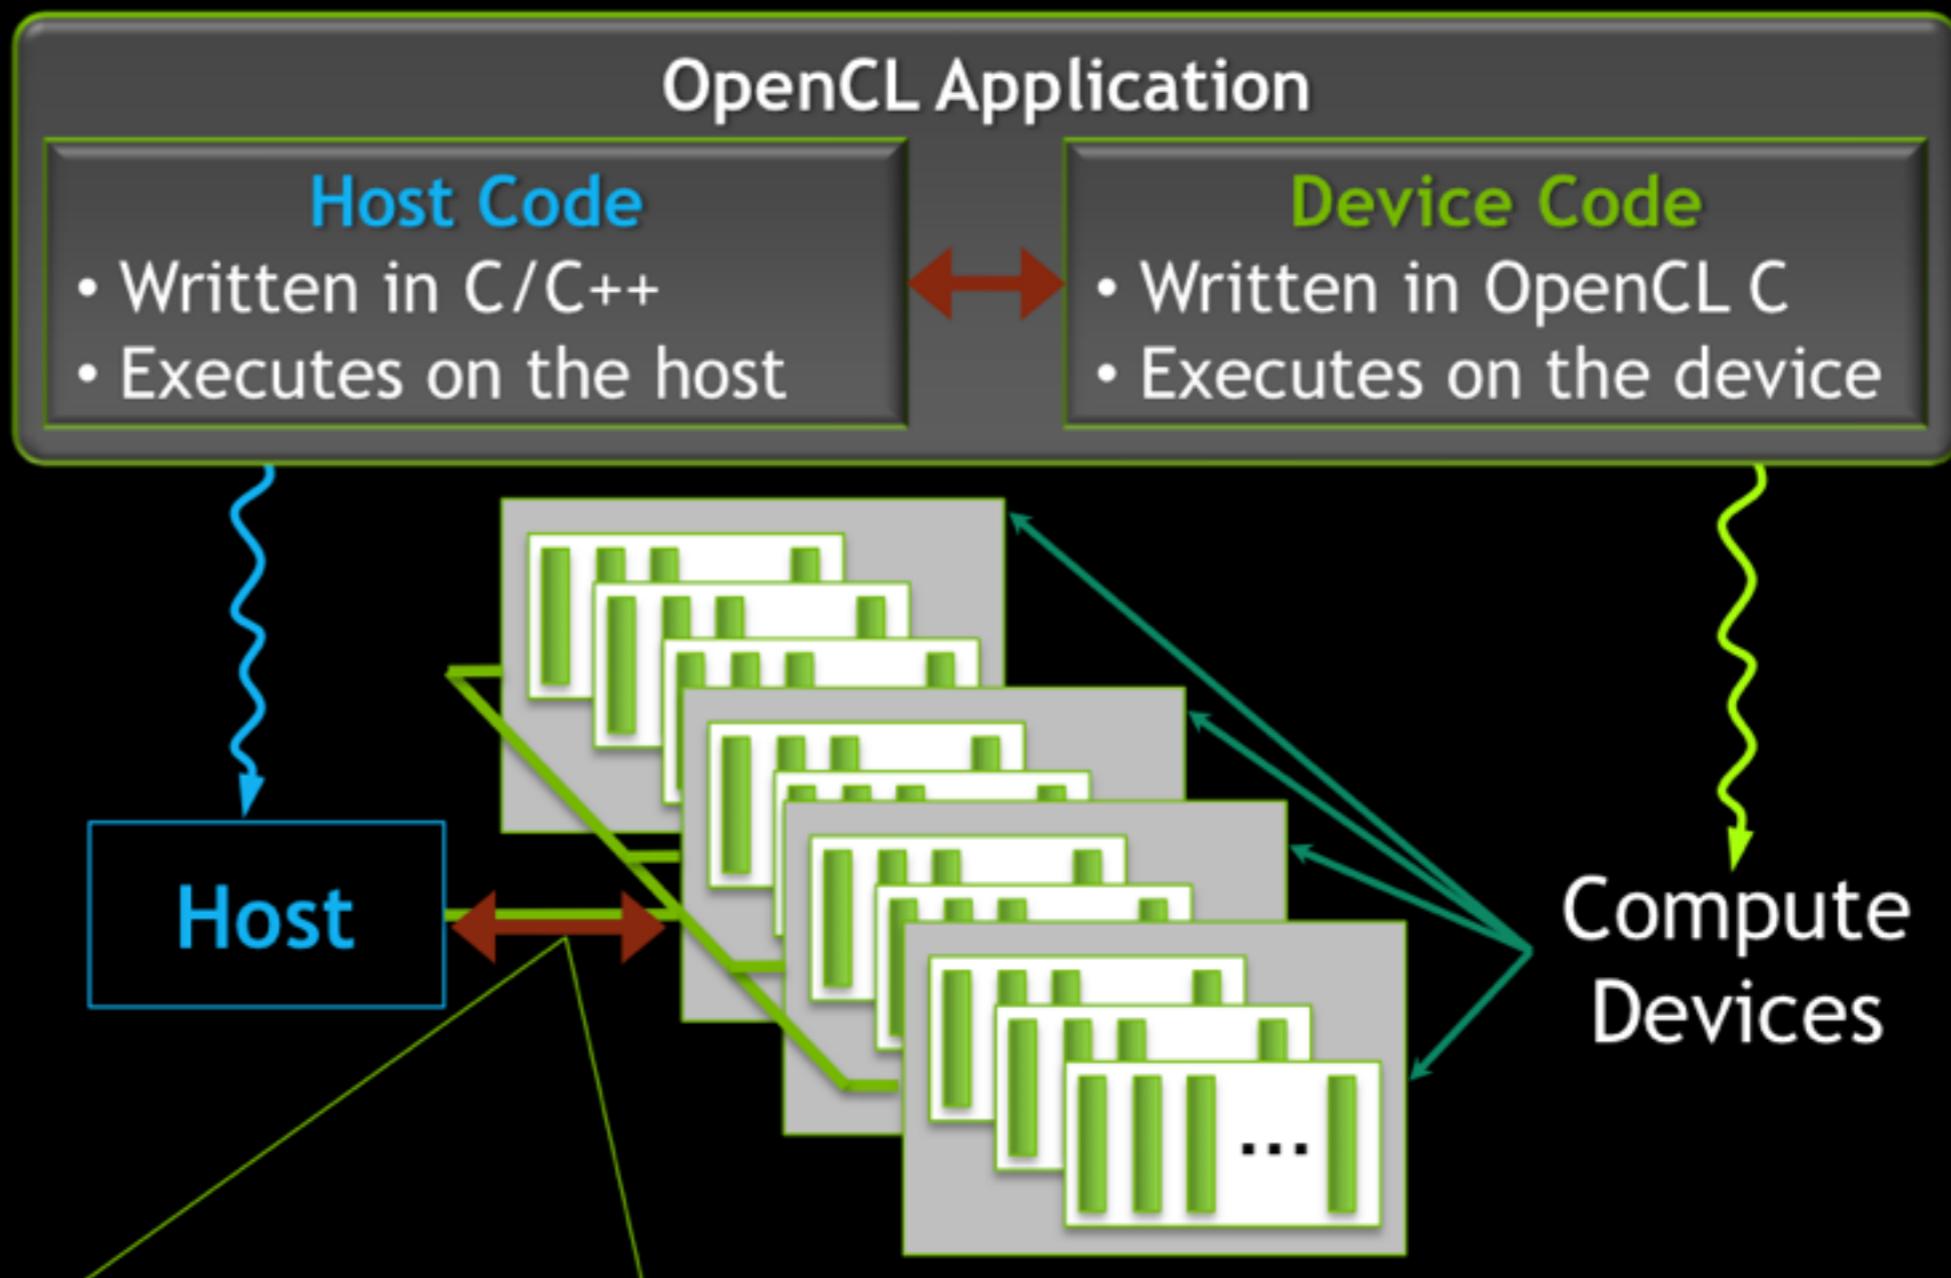
OpenCL



OpenCL

PART III
CHAPTER 1

Anatomy of an OpenCL Application



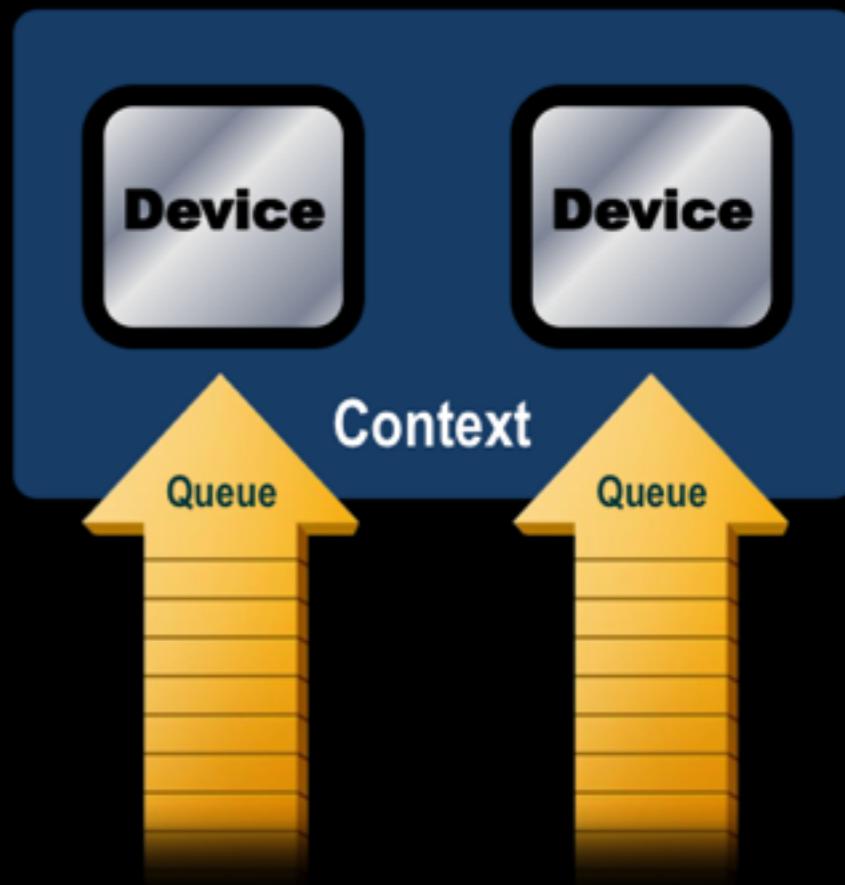
Host code sends commands to the **Devices**:

- ... to transfer data between host memory and device memories
- ... to execute device code

OpenCL Execution Model

The application runs on a **Host** which submits work to the **Devices**

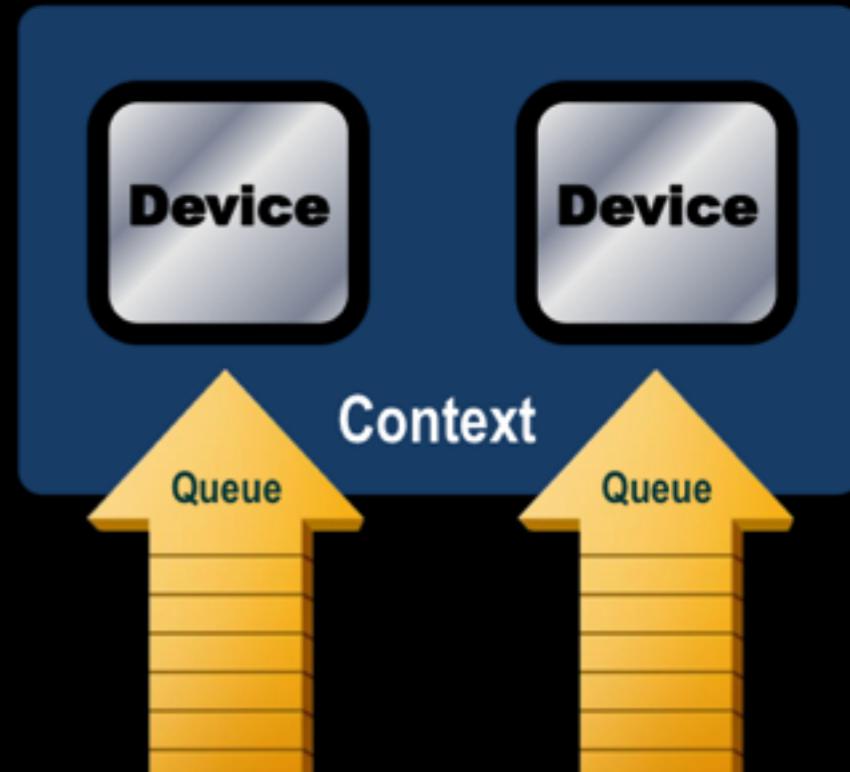
- **Context:** The environment within which work-items execute; includes devices and their memories and command queues
- **Command Queue:** A queue used by the Host application to submit work to a Device (e.g., kernel execution instances)
 - Work is queued in-order, one queue per device
 - Work can be executed in-order *or* out-of-order



OpenCL Execution Model

The application runs on a **Host** which submits work to the **Devices**

- **Context:** The environment within which work-items execute; includes devices and their memories and command queues
- **Command Queue:** A queue used by the Host application to submit work to a Device (e.g., kernel execution instances)
 - Work is queued in-order, one queue per device
 - Work can be executed in-order *or* out-of-order

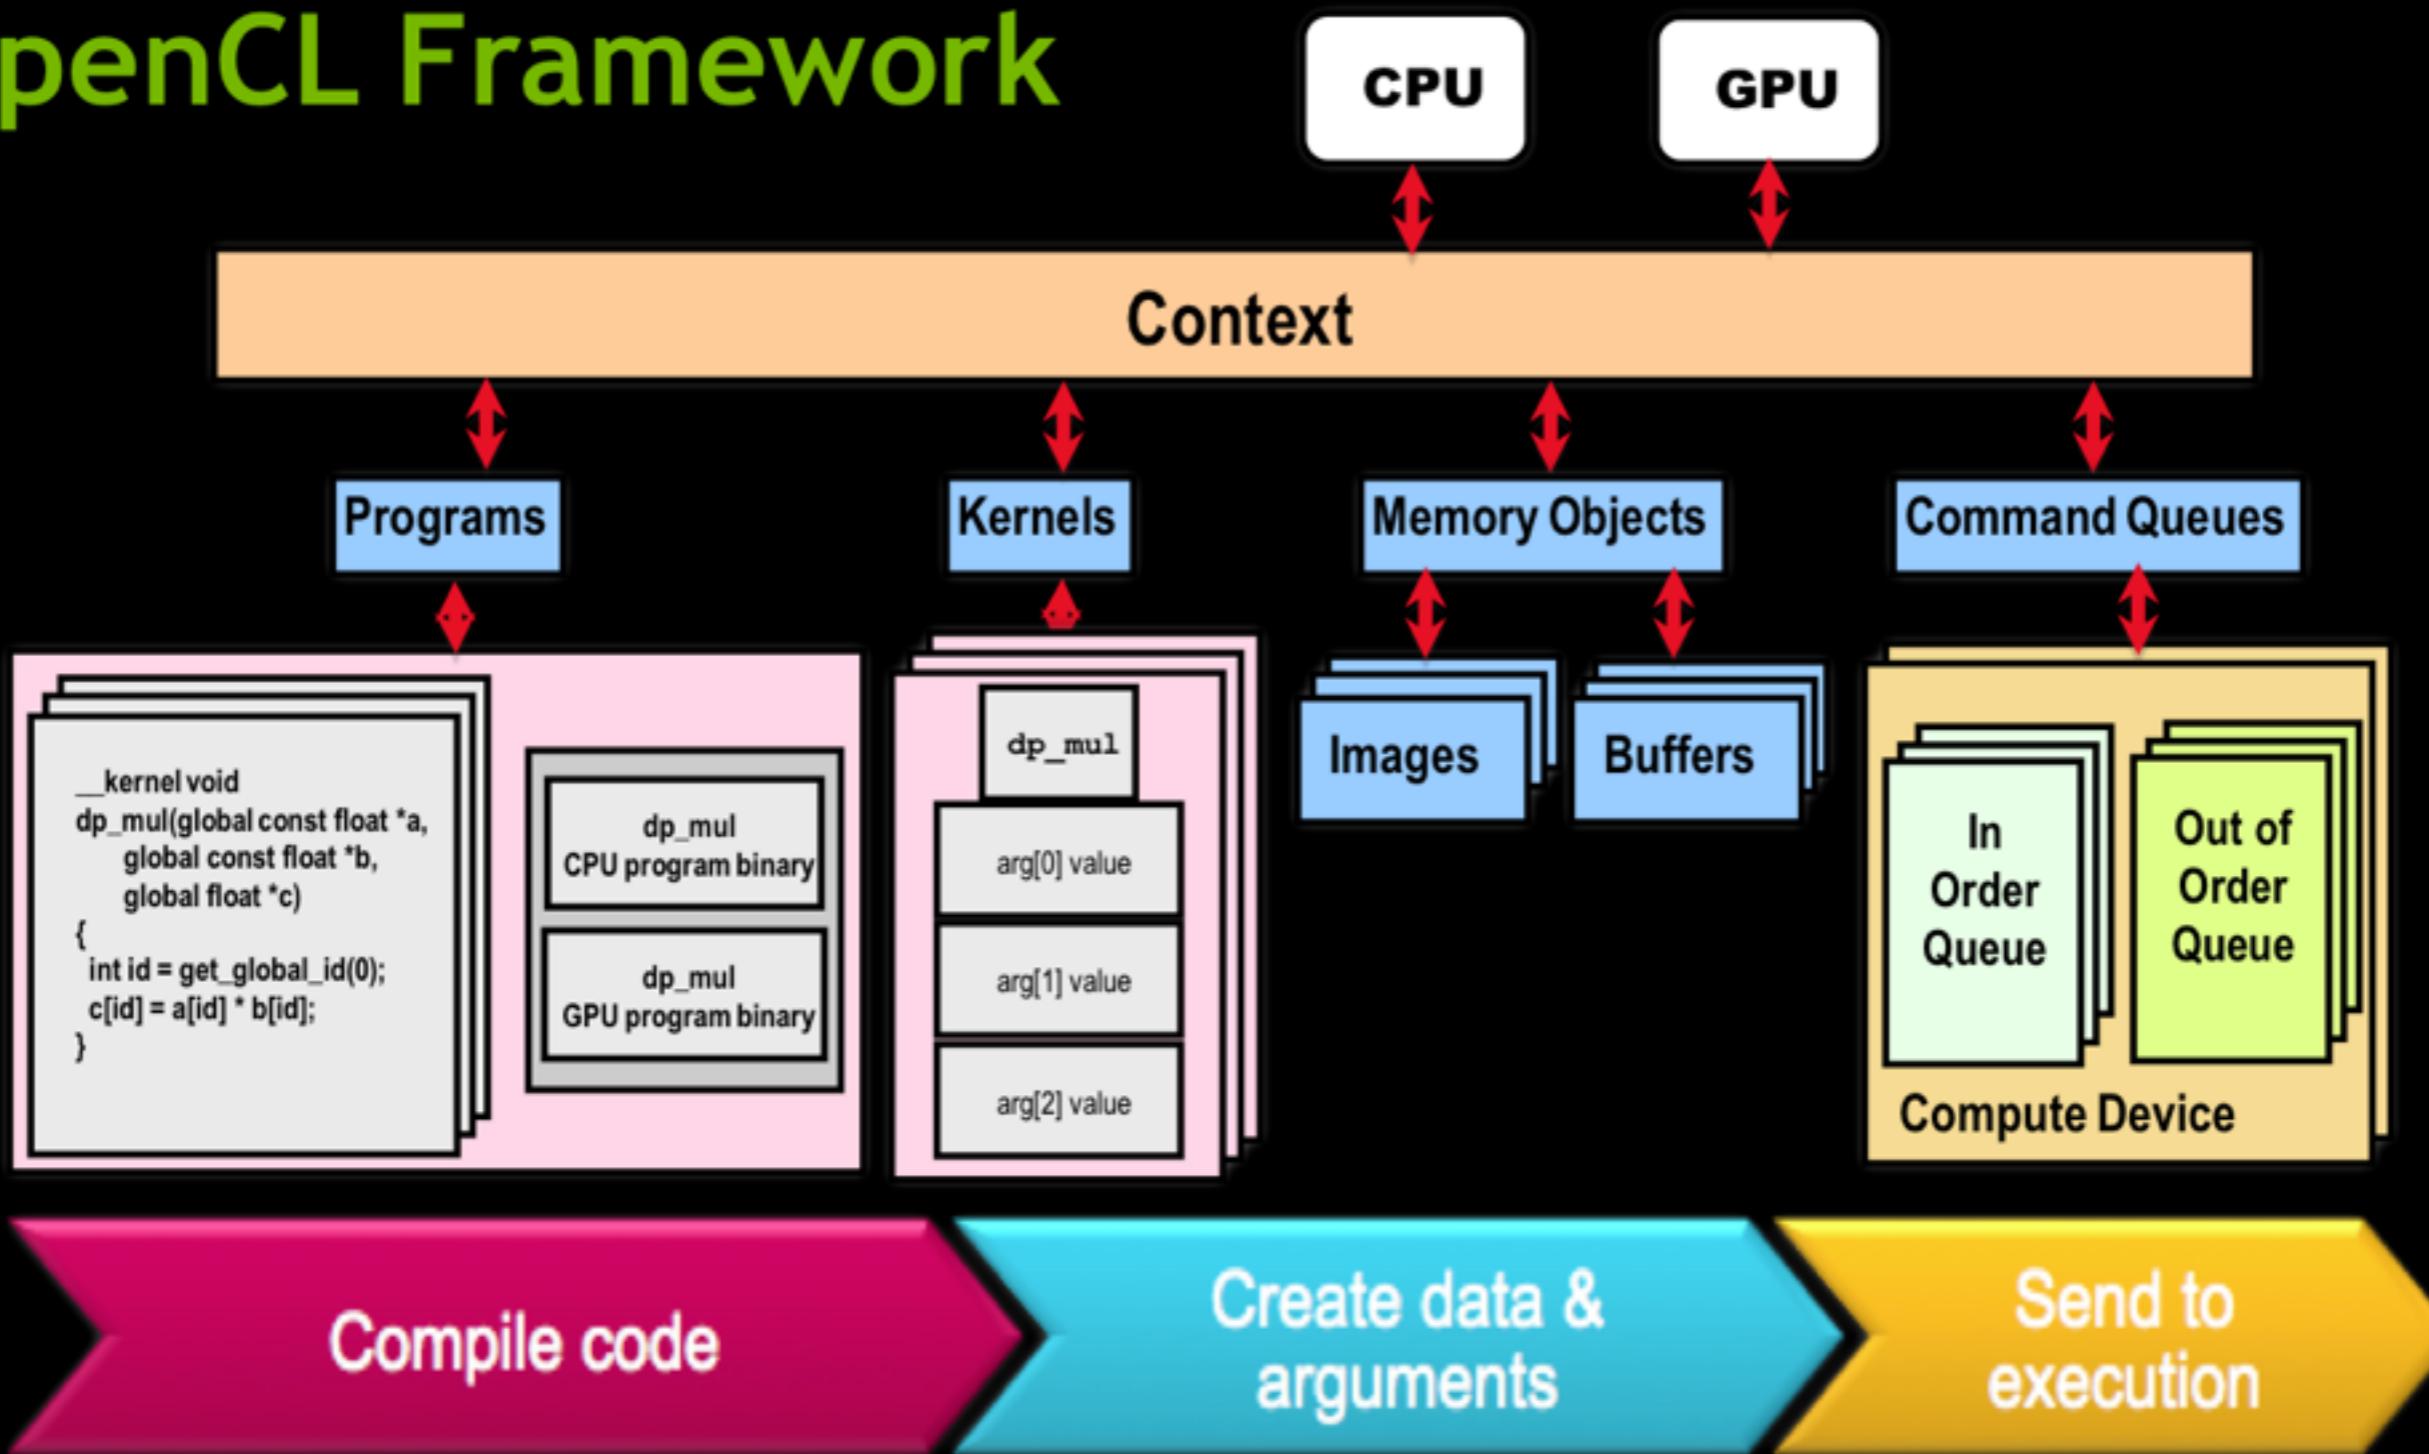


```
// Connect to a compute device
int gpu = 1;
err = clGetDeviceIDs(NULL, gpu ? CL_DEVICE_TYPE_GPU : CL_DEVICE_TYPE_CPU,
                     1, &device_id, NULL);

// Create a compute context
context = clCreateContext(0, 1, &device_id, NULL, NULL, &err);

// Create a command queue
commands = clCreateCommandQueue(context, device_id, 0, &err);
```

OpenCL Framework



OpenCL Framework: Platform Layer

- Query platform information
 - `clGetPlatformInfo()`: profile, version, vendor, extensions
 - `clGetDeviceIDs()`: list of devices
 - `clDeviceInfo()`: type, capabilities
- Create an OpenCL context for one or more devices

Context = {

- One or more devices
`cl_device_id`
- Memory and device code shared by these devices
`cl_mem` `cl_program`
- Command queues to send commands to these devices
`cl_command_queue`

KERNEL

Decompose task into *work-items*

- Define N-dimensional computation domain
- Execute a *kernel* at each point in computation domain

Traditional loop as a function in C

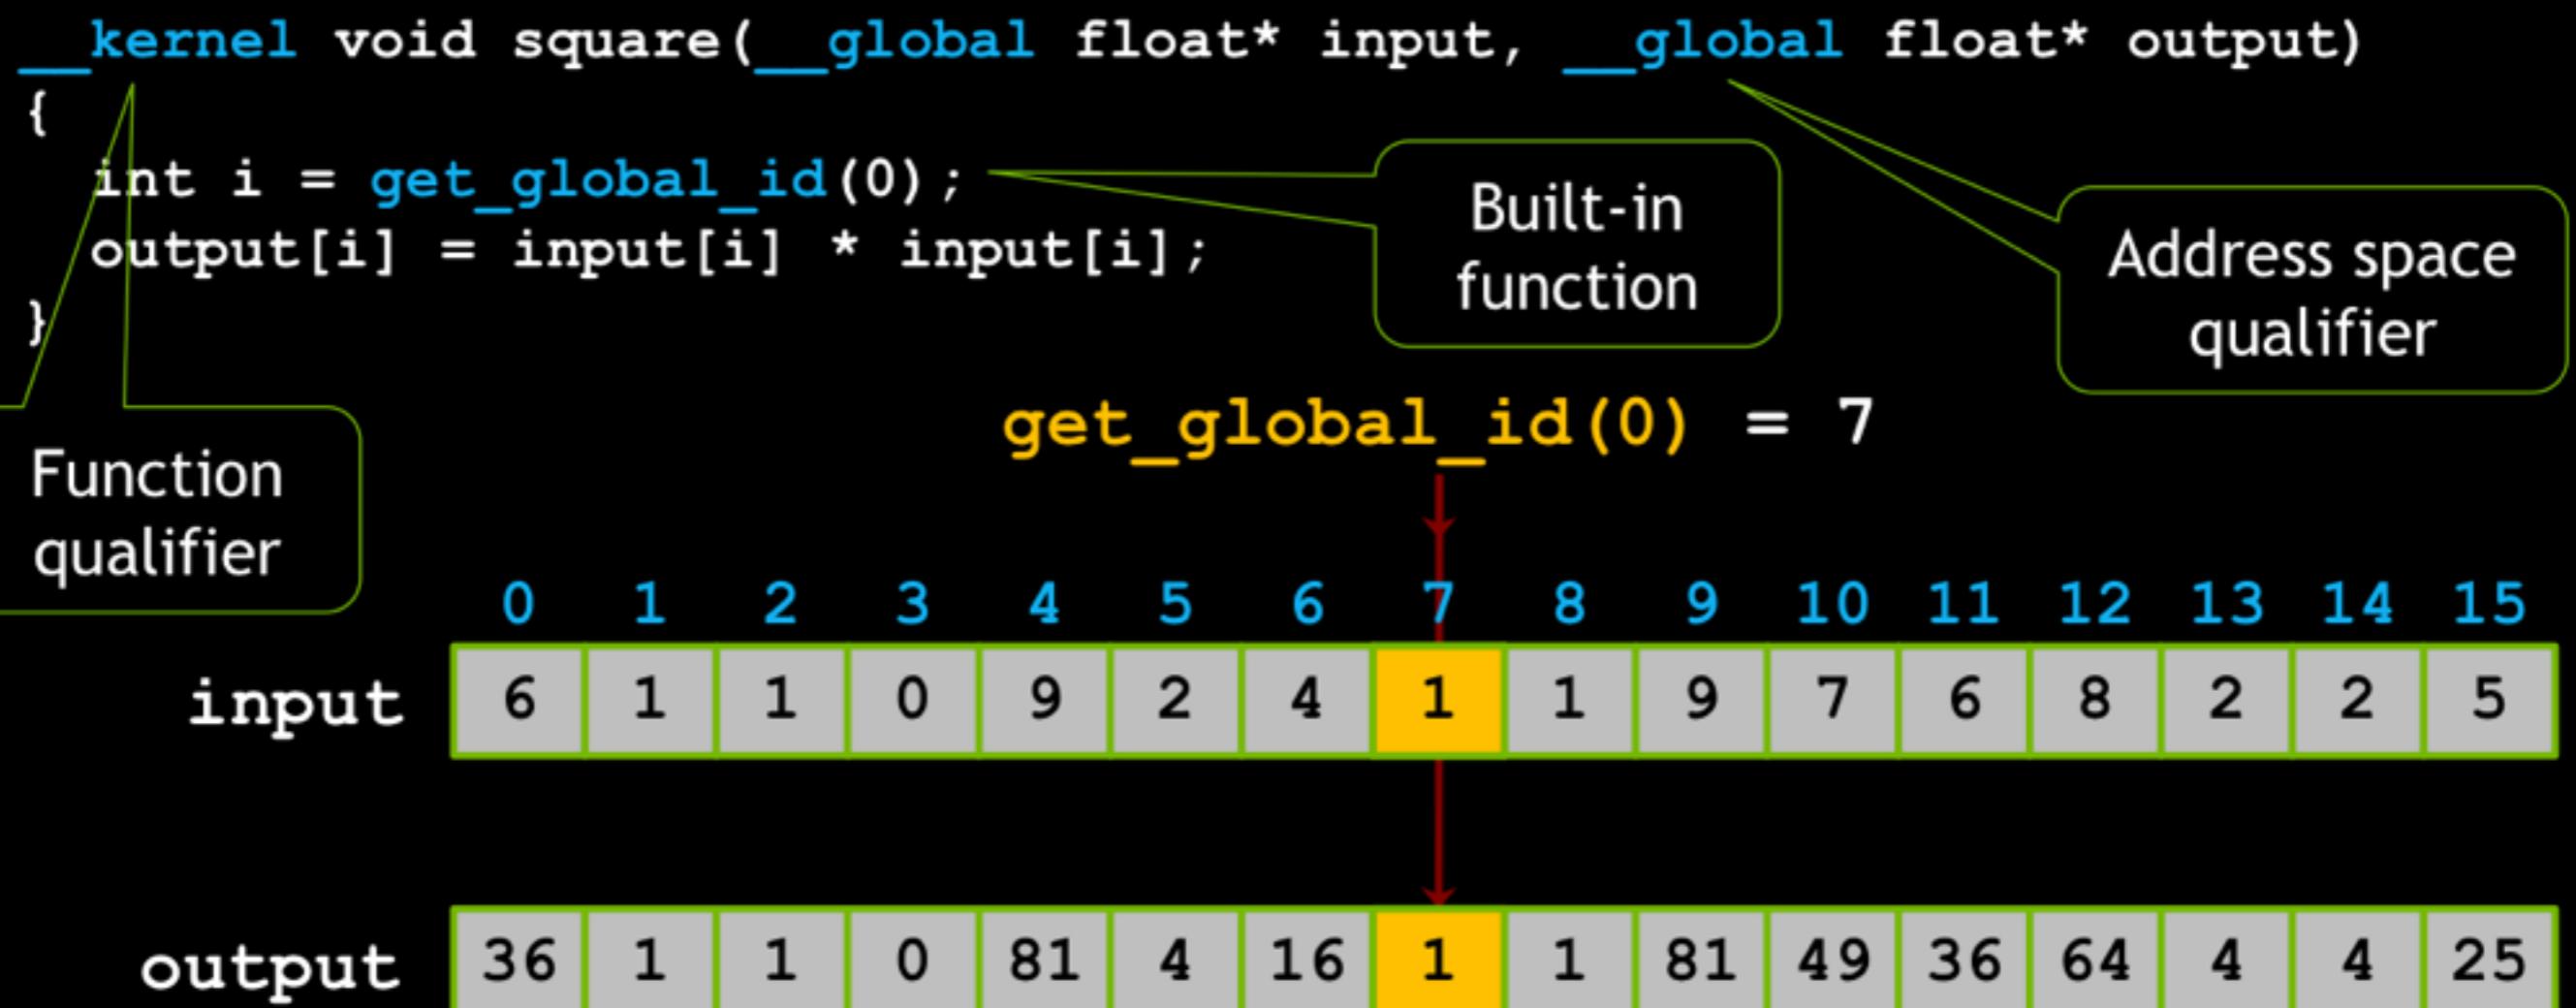
```
void  
trad_mul(int n,  
         const float *a,  
         const float *b,  
         float *c)  
{  
    int i;  
    for (i=0; i<n; i++)  
        c[i] = a[i] * b[i];  
}
```

OpenCL C kernel

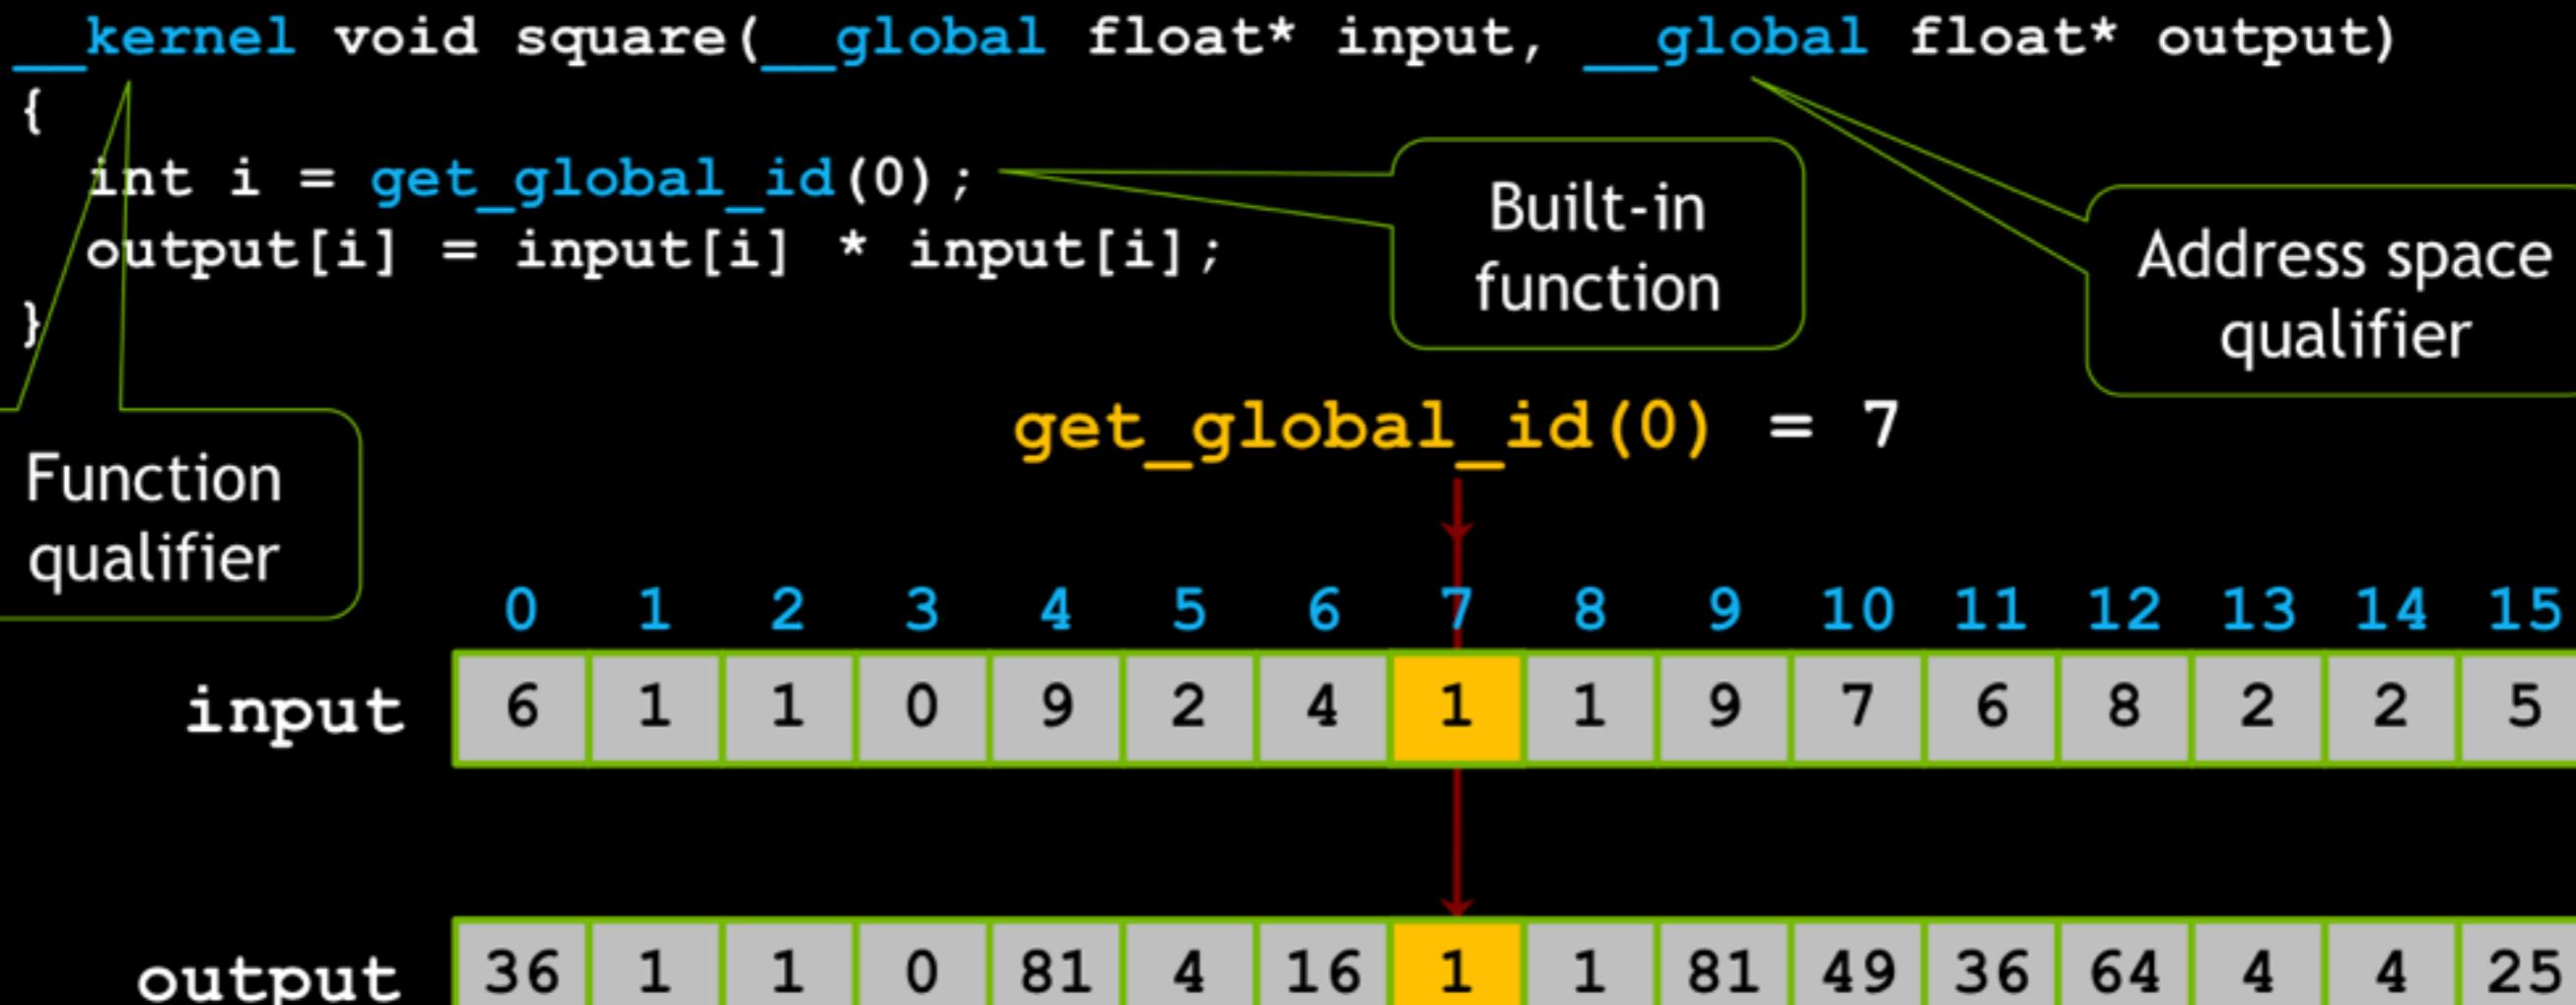
```
__kernel void  
dp_mul(__global const float *a,  
       __global const float *b,  
       __global float *c)  
{  
    int id = get_global_id(0);  
  
    c[id] = a[id] * b[id];  
} // execute over n "work items"
```



- A *kernel* is a function executed for each work-item



- A *kernel* is a function executed for each work-item



```
// Create the compute program from the source buffer
program = clCreateProgramWithSource(context, 1, (const char **) &KernelSource, NULL, &err);

// Build the program executable
err = clBuildProgram(program, 0, NULL, NULL, NULL, NULL);

// Create the compute kernel in the program we wish to run
kernel = clCreateKernel(program, "sqrt", &err);
```

WRITE AND READ DATA ON GPU

```
// Create the input and output arrays in device memory  
for our calculation  
input = clCreateBuffer(context, CL_MEM_READ_ONLY,  
                      sizeof(float) * data_size,  
                      NULL, NULL);  
output = clCreateBuffer(context, CL_MEM_WRITE_ONLY,  
                      sizeof(float) * data_size,  
                      NULL, NULL);  
  
// Write our data set into the input array in device  
memory  
err = clEnqueueWriteBuffer(commands, input, CL_TRUE, 0,  
                           sizeof(float) * data_size, data, 0, NULL, NULL);
```

WRITE AND READ DATA ON GPU

```
// Create the input and output arrays in device memory  
for our calculation  
input = clCreateBuffer(context, CL_MEM_READ_ONLY,  
                      sizeof(float) * data_size,  
                      NULL, NULL);  
output = clCreateBuffer(context, CL_MEM_WRITE_ONLY,  
                      sizeof(float) * data_size,  
                      NULL, NULL);  
  
// Write our data set into the input array in device  
memory  
err = clEnqueueWriteBuffer(commands, input, CL_TRUE, 0,  
                           sizeof(float) * data_size, data, 0, NULL, NULL);
```

... run computations here ...

WRITE AND READ DATA ON GPU

```
// Create the input and output arrays in device memory  
for our calculation  
input = clCreateBuffer(context, CL_MEM_READ_ONLY,  
                      sizeof(float) * data_size,  
                      NULL, NULL);  
output = clCreateBuffer(context, CL_MEM_WRITE_ONLY,  
                      sizeof(float) * data_size,  
                      NULL, NULL);  
  
// Write our data set into the input array in device  
memory  
err = clEnqueueWriteBuffer(commands, input, CL_TRUE, 0,  
                           sizeof(float) * data_size, data, 0, NULL, NULL);  
  
... run computations here ...  
  
// Read back the results from the device  
err = clEnqueueReadBuffer( commands, output, CL_TRUE, 0,  
                           sizeof(float) * data_size, gpu_results, 0, NULL, NULL );
```

THE COMPUTATION

```
// Set the arguments to our compute kernel
err = 0;
err |= clSetKernelArg(kernel, 0, sizeof(cl_mem), &input);
err |= clSetKernelArg(kernel, 1, sizeof(cl_mem), &output);
err |= clSetKernelArg(kernel, 2, sizeof(unsigned long), &data_size);
```

THE COMPUTATION

```
// Set the arguments to our compute kernel
err = 0;
err |= clSetKernelArg(kernel, 0, sizeof(cl_mem), &input);
err |= clSetKernelArg(kernel, 1, sizeof(cl_mem), &output);
err |= clSetKernelArg(kernel, 2, sizeof(unsigned long), &data_size);
```

```
kernel void square( global float* input,
                      global float* output,
                      const unsigned long data_size)
{
    int i = get_global_id(0);
    if(i < data_size)
        output[i] = sqrt(input[i]);
}
```

THE COMPUTATION

```
// Set the arguments to our compute kernel
err = 0;
err |= clSetKernelArg(kernel, 0, sizeof(cl_mem), &input);
err |= clSetKernelArg(kernel, 1, sizeof(cl_mem), &output);
err |= clSetKernelArg(kernel, 2, sizeof(unsigned long), &data_size);

// Get the maximum work group size for executing
// the kernel on the device
err = clGetKernelWorkGroupInfo(kernel, device_id,
                                CL_KERNEL_WORK_GROUP_SIZE,
                                sizeof(local), &local, NULL);
```

```
kernel void square( __global float* input,
                    __global float* output,
                    const unsigned long data_size)
{
    int i = get_global_id(0);
    if(i < data_size)
        output[i] = sqrt(input[i]);
}
```

THE COMPUTATION

```
// Set the arguments to our compute kernel
err = 0;
err = clSetKernelArg(kernel, 0, sizeof(cl_mem), &input);
err |= clSetKernelArg(kernel, 1, sizeof(cl_mem), &output);
err |= clSetKernelArg(kernel, 2, sizeof(unsigned long), &data_size);

// Get the maximum work group size for executing
// the kernel on the device
err = clGetKernelWorkGroupInfo(kernel, device_id,
                               CL_KERNEL_WORK_GROUP_SIZE,
                               sizeof(local), &local, NULL);

// Execute the kernel over the entire range of our 1D input data set
// using the maximum number of work group items for this device
global = data_size;
err = clEnqueueNDRangeKernel(commands, kernel, 1, NULL,
                             &global, &local, 0, NULL, NULL);
```

THE COMPUTATION

```
// Set the arguments to our compute kernel
err = 0;
err = clSetKernelArg(kernel, 0, sizeof(cl_mem), &input);
err |= clSetKernelArg(kernel, 1, sizeof(cl_mem), &output);
err |= clSetKernelArg(kernel, 2, sizeof(unsigned long), &data_size);

// Get the maximum work group size for executing
// the kernel on the device
err = clGetKernelWorkGroupInfo(kernel, device_id,
                               CL_KERNEL_WORK_GROUP_SIZE,
                               sizeof(local), &local, NULL);

// Execute the kernel over the entire range of our 1D input data set
// using the maximum number of work group items for this device
global = data_size;
err = clEnqueueNDRangeKernel(commands, kernel, 1, NULL,
                             &global, &local, 0, NULL, NULL);

// Wait for the command commands to get serviced before
// reading back results
clFinish(commands);
```

DEMO

The screenshot shows the Xcode IDE interface for an OpenCL project named "OpenCL". The project structure on the left includes "OpenCL", "Frameworks" (with "OpenCL.framework"), "Sources" (containing "opencl.c"), and "Products" (containing "OpenCL"). The main editor window displays the "opencl.c" source code, which contains C code for initializing an OpenCL context, command queue, program, and kernel, and creating input and output buffers. The code uses the CL API functions like clGetDeviceIDs, clCreateContext, clCreateCommandQueue, clCreateProgramWithSource, clBuildProgram, clCreateKernel, and clCreateBuffer. The bottom right panel shows the build output, indicating a CPU time of 1.236780 and a GPU time of 0.014359, with a status message: "Program ended with exit code: 0".

```
// Connect to a compute device
int gpu = 1;
err = clGetDeviceIDs(NULL, gpu ? CL_DEVICE_TYPE_GPU : CL_DEVICE_TYPE_CPU, 1, &device_id, NULL);

// Create a compute context
context = clCreateContext(0, 1, &device_id, NULL, NULL, &err);

// Create a command queue
commands = clCreateCommandQueue(context, device_id, 0, &err);

// Create the compute program from the source buffer
program = clCreateProgramWithSource(context, 1, (const char **) &KernelSource, NULL,
    &err);

// Build the program executable
err = clBuildProgram(program, 0, NULL, NULL, NULL, NULL);

// Create the compute kernel in the program we wish to run
kernel = clCreateKernel(program, "sqrt", &err);

// Create the input and output arrays in device memory for our calculation
input = clCreateBuffer(context, CL_MEM_READ_ONLY, sizeof(float) * data_size, NULL,
    NULL);
output = clCreateBuffer(context, CL_MEM_WRITE_ONLY, sizeof(float) * data_size, NULL,
    NULL);
```

CPU time 1.236780
GPU time 0.014359
Program ended with exit code: 0

Open, study and run the project from the code/OpenCL



PART III
CHAPTER 2



CUDA PROGRAMMING MODEL

- CPU is called “*host*”
 - Move data CPU <-> GPU memory `cudaMemcpy`
 - Allocate memory `cudaMalloc`
 - Launch *kernels* on GPU
- GPU is called “*device*”



CUDA PROGRAMMING MODEL

- CPU is called “host”
 - Move data CPU <-> GPU memory `cudaMemcpy`
 - Allocate memory `cudaMalloc`
 - Launch *kernels* on GPU
- GPU is called “device”

Typical CUDA program

1. CPU allocates memory on GPU
2. CPU copies data to GPU memory
3. CPU launches kernels on GPU (process the data)
4. CPU copies results back to CPU memory



CUDA PROGRAMMING MODEL

- CPU is called “host”
 - Move data CPU <-> GPU memory `cudaMemcpy`
 - Allocate memory `cudaMalloc`
 - Launch *kernel*s on GPU
- GPU is called “device”

Very similar to the logic of OpenCL

Typical CUDA workflow

1. CPU allocates memory on GPU
2. CPU copies data to GPU memory
3. CPU launches kernels on GPU (process the data)
4. CPU copies results back to CPU memory

EXAMPLE

```
#include <stdio.h>

// Kernel
__global__ void cube(float * d_out, float * d_in){
    int id = threadIdx.x;
    d_out[id] = d_in[id] * d_in[id] * d_in[id];
}

int main(int argc, char ** argv) {
    const int ARRAY_SIZE = 96;
    const int ARRAY_BYTES = ARRAY_SIZE * sizeof(float);

    // generate the input array on the host
    float h_in[ARRAY_SIZE];
    for (int i = 0; i < ARRAY_SIZE; i++) {
        h_in[i] = float(i);
    }
    float h_out[ARRAY_SIZE];

    // declare GPU memory pointers
    float * d_in;
    float * d_out;

    // allocate GPU memory
    cudaMalloc((void**) &d_in, ARRAY_BYTES);
    cudaMalloc((void**) &d_out, ARRAY_BYTES);
```

EXAMPLE

```
#include <stdio.h>

// Kernel
__global__ void cube(float * d_out, float * d_in){
    int id = threadIdx.x;
    d_out[id] = d_in[id] * d_in[id] * d_in[id];
}

int main(int argc, char ** argv) {
    const int ARRAY_SIZE = 96;
    const int ARRAY_BYTES = ARRAY_SIZE * sizeof(float);

    // generate the input array on the host
    float h_in[ARRAY_SIZE];
    for (int i = 0; i < ARRAY_SIZE; i++) {
        h_in[i] = float(i);
    }
    float h_out[ARRAY_SIZE];

    // declare GPU memory pointers
    float * d_in;
    float * d_out;

    // allocate GPU memory
    cudaMalloc((void**) &d_in, ARRAY_BYTES);
    cudaMalloc((void**) &d_out, ARRAY_BYTES);
```

EXAMPLE

```
#include <stdio.h>

// Kernel
__global__ void cube(float * d_out, float * d_in){
    int id = threadIdx.x;
    d_out[id] = d_in[id] * d_in[id] * d_in[id];
}

int main(int argc, char ** argv) {
    const int ARRAY_SIZE = 96;
    const int ARRAY_BYTES = ARRAY_SIZE * sizeof(float);

    // generate the input array on the host
    float h_in[ARRAY_SIZE];
    for (int i = 0; i < ARRAY_SIZE; i++) {
        h_in[i] = float(i);
    }
    float h_out[ARRAY_SIZE];

    // declare GPU memory pointers
    float * d_in;
    float * d_out;

    // allocate GPU memory
    cudaMalloc((void**) &d_in, ARRAY_BYTES);
    cudaMalloc((void**) &d_out, ARRAY_BYTES);
```

EXAMPLE

```
// transfer the array to the GPU
cudaMemcpy(d_in, h_in, ARRAY_BYTES, cudaMemcpyHostToDevice);

// launch the kernel
cube<<<1, ARRAY_SIZE>>>(d_out, d_in);

// copy back the result array to the CPU
cudaMemcpy(h_out, d_out, ARRAY_BYTES, cudaMemcpyDeviceToHost);

// print out the resulting array
for (int i =0; i < ARRAY_SIZE; i++) {
    printf("%f", h_out[i]);
}

cudaFree(d_in);
cudaFree(d_out);

return 0;
}
```

EXAMPLE

```
// transfer the array to the GPU
cudaMemcpy(d_in, h_in, ARRAY_BYTES, cudaMemcpyHostToDevice);

// launch the kernel
cube<<<1, ARRAY_SIZE>>>(d_out, d_in);

// copy back the result array to the CPU
cudaMemcpy(h_out, d_out, ARRAY_BYTES, cudaMemcpyDeviceToHost);

// print out the resulting array
for (int i =0; i < ARRAY_SIZE; i++) {
    printf("%.f", h_out[i]);
}

cudaFree(d_in);
cudaFree(d_out);

return 0;
}
```

EXAMPLE

```
// transfer the array to the GPU
cudaMemcpy(d_in, h_in, ARRAY_BYTES, cudaMemcpyHostToDevice);

// launch the kernel
cube<<<1, ARRAY_SIZE>>>(d_out, d_in);

// copy back the result array to the CPU
cudaMemcpy(h_out, d_out, ARRAY_BYTES, cudaMemcpyDeviceToHost);

// print out the resulting array
for (int i =0; i < ARRAY_SIZE; i++) {
    printf("%f", h_out[i]);
}

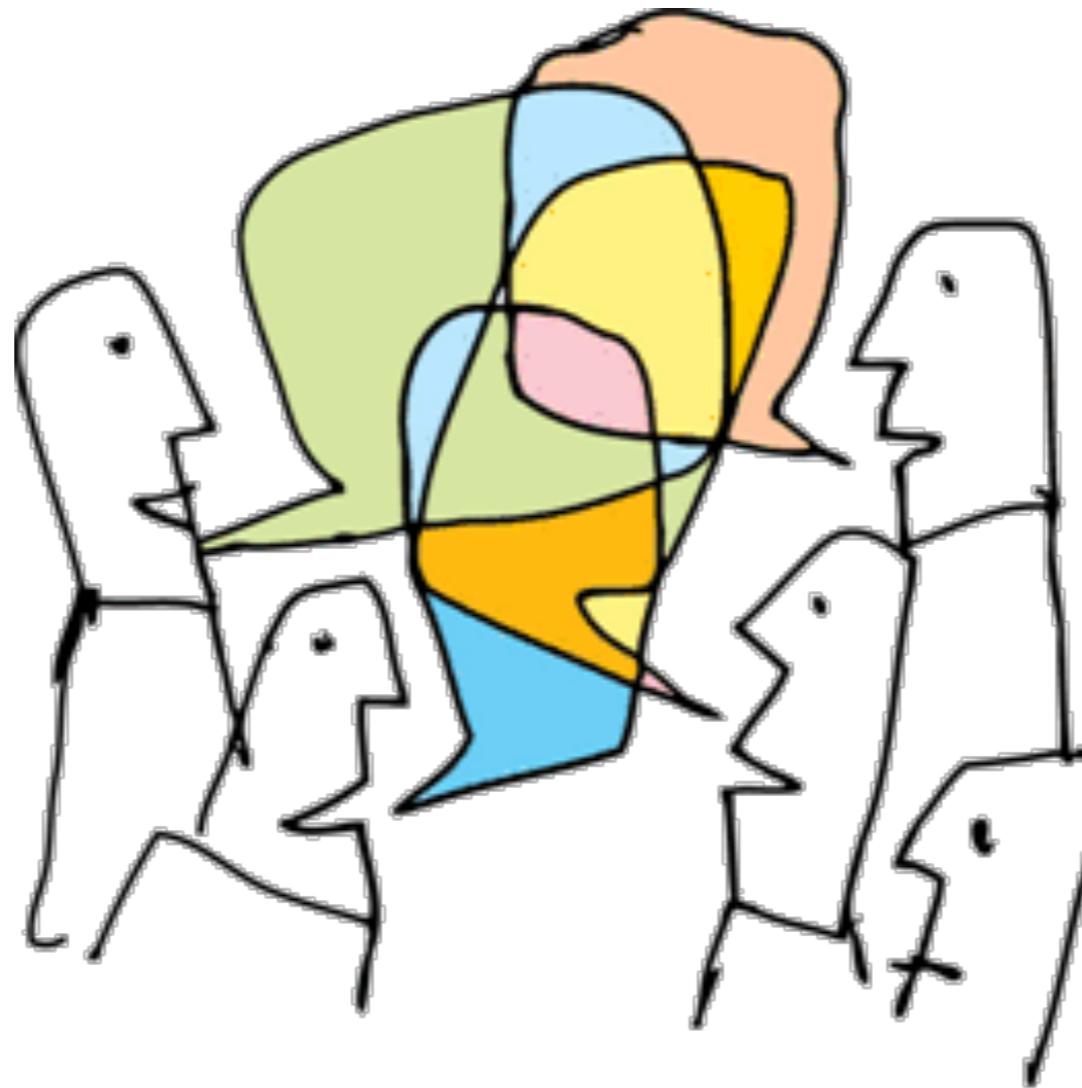
cudaFree(d_in);
cudaFree(d_out);

return 0;
}
```

EXAMPLE

```
// transfer the array to the GPU  
cudaMemcpy(d_in, h_in, ARRAY_BYTES, cudaMemcpyHostToDevice);  
  
// launch the kernel  
cube<<<1, ARRAY_SIZE>>>(d_out, d_in);
```

```
// copy  
cudaMem  
  
// prin  
for (in  
    pri  
}  
    256.000000  
    400.000000  
    576.000000  
    784.000000  
    1024.000000  
    1296.000000  
    1600.000000  
    1936.000000  
    2304.000000  
    2704.000000  
    3136.000000  
    3600.000000  
}  
return  
}  
  
$ nvcc -o square square.cu  
$ ./square  
0.000000      1.000000      4.000000      9.000000  
16.000000     25.000000     36.000000     49.000000  
64.000000     81.000000     100.000000    121.000000  
144.000000    169.000000    196.000000    225.000000  
256.000000    289.000000    324.000000    361.000000  
400.000000    441.000000    484.000000    529.000000  
576.000000    625.000000    676.000000    729.000000  
784.000000    841.000000    900.000000    961.000000  
1024.000000   1089.000000   1156.000000   1225.000000  
1296.000000   1369.000000   1444.000000   1521.000000  
1600.000000   1681.000000   1764.000000   1849.000000  
1936.000000   2025.000000   2116.000000   2209.000000  
2304.000000   2401.000000   2500.000000   2601.000000  
2704.000000   2809.000000   2916.000000   3025.000000  
3136.000000   3249.000000   3364.000000   3481.000000  
3600.000000   3721.000000   3844.000000   3969.000000  
$
```



PART IV
“DISCUSSION”

LINKS

- Code repository for this presentation
 - <https://github.com/kuz/Introduction-to-GPU-Computing>
 - Feel free to leave feature requests there, ask questions, etc.
- OpenCL tutorials & presentations
 - <http://streamcomputing.eu/knowledge/for-developers/tutorials/>
 - <http://opencl.codeplex.com/wikipage?title=OpenCL%20Tutorials%20-%201&referringTitle=OpenCL%20Tutorials>
 - http://www.cc.gatech.edu/~vetter/keeneland/tutorial-2011-04-14/06-intro_to_opencl.pdf
- Introduction to Parallel Computing @ www.udacity.com
 - <https://www.udacity.com/course/cs344>
- Slides about Compute Shader
 - <http://web.engr.oregonstate.edu/~mjb/cs557/Handouts/compute.shader.1pp.pdf>
- Introduction to Computer Graphics with codebases
 - <https://github.com/konstantint/ComputerGraphics2013>
- GLSL Computing (aka “Old school”)
 - <http://www.computer-graphics.se/gpu-computing/lab1.html>